

Software Design Document: GPU N-Body Integration Toolkit

Prepared by:

Daniel Bagnell

Jason Economou

Rajkumar Jayachandran

Tim McJilton

Gabe Schwartz

Andrew Sherman

Advisor:

Prof. Jeremy Johnson

Stakeholders:

Prof. Steve McMillan

Alfred Whitehead

Contents

1	Introduction	4
1.1	Purpose	4
1.2	Scope	4
1.3	Glossary	4
2	High Level Architecture	5
3	Graphical User Interface	6
3.1	Overview	6
3.2	Network Interaction	6
3.3	Node Management	7
4	Experiment Components	9
4.1	Experiment Object Model	9
4.1.1	Experiment Class	10
4.1.2	Experiment Hierarchy	13
4.2	Particles	14
4.2.1	Particle Class	14
4.3	Modules	16
4.3.1	Module Class	16
4.3.2	Parameter Class	20
4.3.3	Unit Class	22
4.3.4	CompoundUnit Class	23
4.3.5	UnitType Enumeration	26
4.3.6	SIPrefix Enumeration	27
4.3.7	AstrophysicalDomain Enumeration	27
4.3.8	StoppingConditions Enumeration	28
4.4	Diagnostics	28
4.4.1	Object Model	28
4.4.2	Diagnostic Class	29
4.4.3	Experiment Manager Class	29
4.4.4	Condition Class	29
4.5	Logging	29
4.5.1	Object Model	29
4.5.2	Logger Class	29
5	Experiment Generation	30
5.1	Overview	30
6	Command Line Tool	30
6.1	Overview	30
6.2	Using the CLT	31
6.2.1	CLT Flags	31

6.3	Running The Experiment	31
6.3.1	Simulation Overview	31
7	Networking	34
7.1	Overview	34
7.2	Network Software Components	35
7.2.1	Control Instance	35
7.2.2	Node Instance	35
7.3	Object Model	35
7.3.1	Overview	35
7.3.2	Control Instance	36
7.3.3	Node Instance	36
7.3.4	TransmissionThread	36
7.3.5	ReceivingThread	37
7.3.6	ListeningThread	38
7.3.7	Packet Objects	39
8	Data	39
8.1	Experiment Specification	39
8.1.1	Experiment	40
8.1.2	Time	40
8.1.3	Module	40
8.1.4	Param	40
8.1.5	Particle	40
8.1.6	Diagnostic	41
8.1.7	Logger	41
8.1.8	Stopping Condition	41
8.2	Experiment Results Directory Structure	41
8.2.1	Top Level Directory	41
8.2.2	Single Experiment Directory	42
8.2.3	Run Directory	42
8.3	Network Packets	42
8.3.1	Packet Header	42
8.3.2	Status Query Packet	42
8.3.3	Status Response Packet	42
8.3.4	Capability Query Packet	42
8.3.5	Capability Response Packet	43
9	Appendix	43
9.1	Requirements Traceability Table	43

1 Introduction

1.1 Purpose

This document serves as the system design for the GPUT framework to control AMUSE experiments as set forth in the Software Requirements Specification. AMUSE (Astrophysical Multipurpose Software Environment) is a software package and set of legacy codes for performing astrophysics simulations. For more details on the purpose and use cases of AMUSE, please refer to the GPUT SRS. The goal of this document is to provide details guiding the construction of the framework, as well as to provide a reference to the framework's architecture. It contains specifications for the objects used by the system and our initial reference implementation plans.

1.2 Scope

This document describes the implementation for our framework, covering the network layer, the user interface and the core python command line scripts. It covers user interface support classes as they will need to be coded, however the actual UI components such as windows and widgets are not described. Code and objects for these items are generated as part of the Qt Creator¹ GUI development kit. As such, the code is not intended to be human-readable and is subject to change between Qt versions. The GPUT SRS document contains the reference prototype which serves as the design for the GUI widgets. The data models including experiment file layouts and network packet details are also covered.

1.3 Glossary

For terms not in this glossary, please refer to the SRS.

Cluster - A group of networked computing nodes available to perform some task.

Computer Node - A computer that is a member of a cluster.

IPC Channel - Any object or memory space used for Inter-Process Communication (IPC). Examples include UNIX Pipes, BSD/Winsock Sockets and UNIX Shared Memory (SHM).

Node Health Statistics - Information about the current state of a node including properties such as CPU usage and memory usage.

¹<http://qt.nokia.com/products/developer-tools/>

2 High Level Architecture

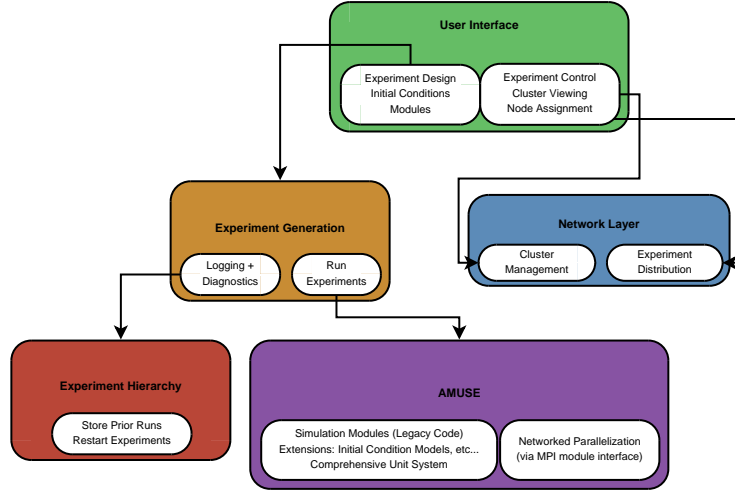


Figure 2.1: Architecture Diagram

Experiments are designed and controlled from the top-level user interface layer. “Experiment” here includes all logging, diagnostics, and module distribution.

An experiment specification file is generated from the user’s input to the interface. This file guides the code generation to produce code that will run the experiment.

This generated code will set up initial conditions of the system such as initial position, velocity etc..., load and initialize the modules selected by the user, and run the experiment. Generated code will control the parallel execution of modules if this feature is enabled.

Custom logging and diagnostic code can be added to the generated code, and the code comes with default diagnostics that are sufficient to reconstruct a crashed/stopped experiment from scratch. Diagnostic/logging output is stored in the experiment hierarchy which is an opaque interface to a list of past runs of an experiment. Storage implementations might include a database, directory structure, or other persistence mechanism.

The GUI will rely on the network to query the status of nodes in the cluster. Any machine on the local network that is running the node management code will be visible to the GUI. From this status view, the user can assign modules in the experiment to run on a specific machine in the cluster or view the running modules on a node.

Modules, units and parallelization are all facilitated through AMUSE library calls. AMUSE contains a complete system of units able to store and represent quantities in a human-readable fashion. AMUSE provides a common interface to “legacy” modules, expert implementations of astrophysics simulation codes. This interface is

centered around the “`evolve_model`” call which advances the module’s simulation one step into the future. Parallelization is realized via the MPI interface enforced in AMUSE’s legacy module framework. Users may specify a list of nodes as an argument to the “`mpirun`” command, and any subsequent calls to module methods will be automatically distributed across the listed nodes.

3 Graphical User Interface

3.1 Overview

The user interface allows the user to create and manage experiments as described in the SRS. Some of the interface code is generated from the Qt Designer GUI development tool, however there are support classes required for the GUI to interact with the network and experiment hierarchy. These classes are detailed below.

The interface allows the user to load experiments from a variety of sources as defined in Section 4.1.2.

3.2 Network Interaction

The network interaction classes allow the UI to send and receive messages through a control instance (Section 7.2.1). The control instance can be either running locally or on a remote machine. The UI uses this link to query the cluster for status updates and control running experiments.

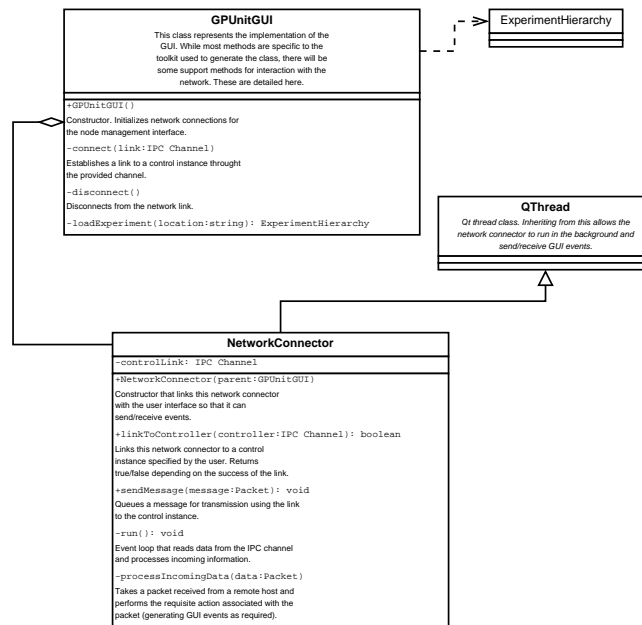


Figure 3.1: GUI Network Connection Class

3.3 Node Management

Nodes in the cluster, detected through the network connector, may be viewed in the node view window. (Satisfies requirements 2.8.1.X)

From the node window, the user can assign experiments to individual nodes. (Satisfies requirement 2.8.2.1)

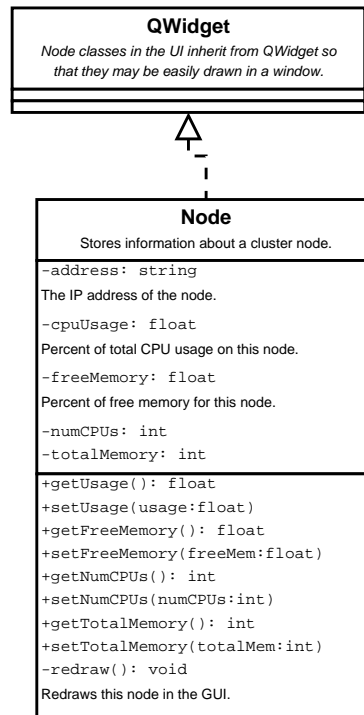


Figure 3.2: Node Class

Figure 3.3 shows an example of an interaction with the network generated by viewing node statuses.

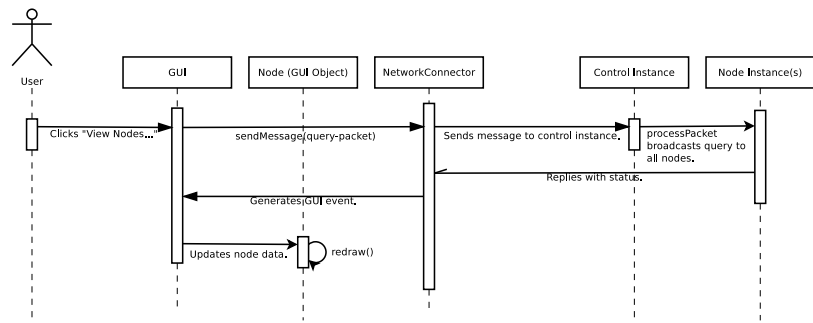


Figure 3.3: Network Interaction

4 Experiment Components

4.1 Experiment Object Model

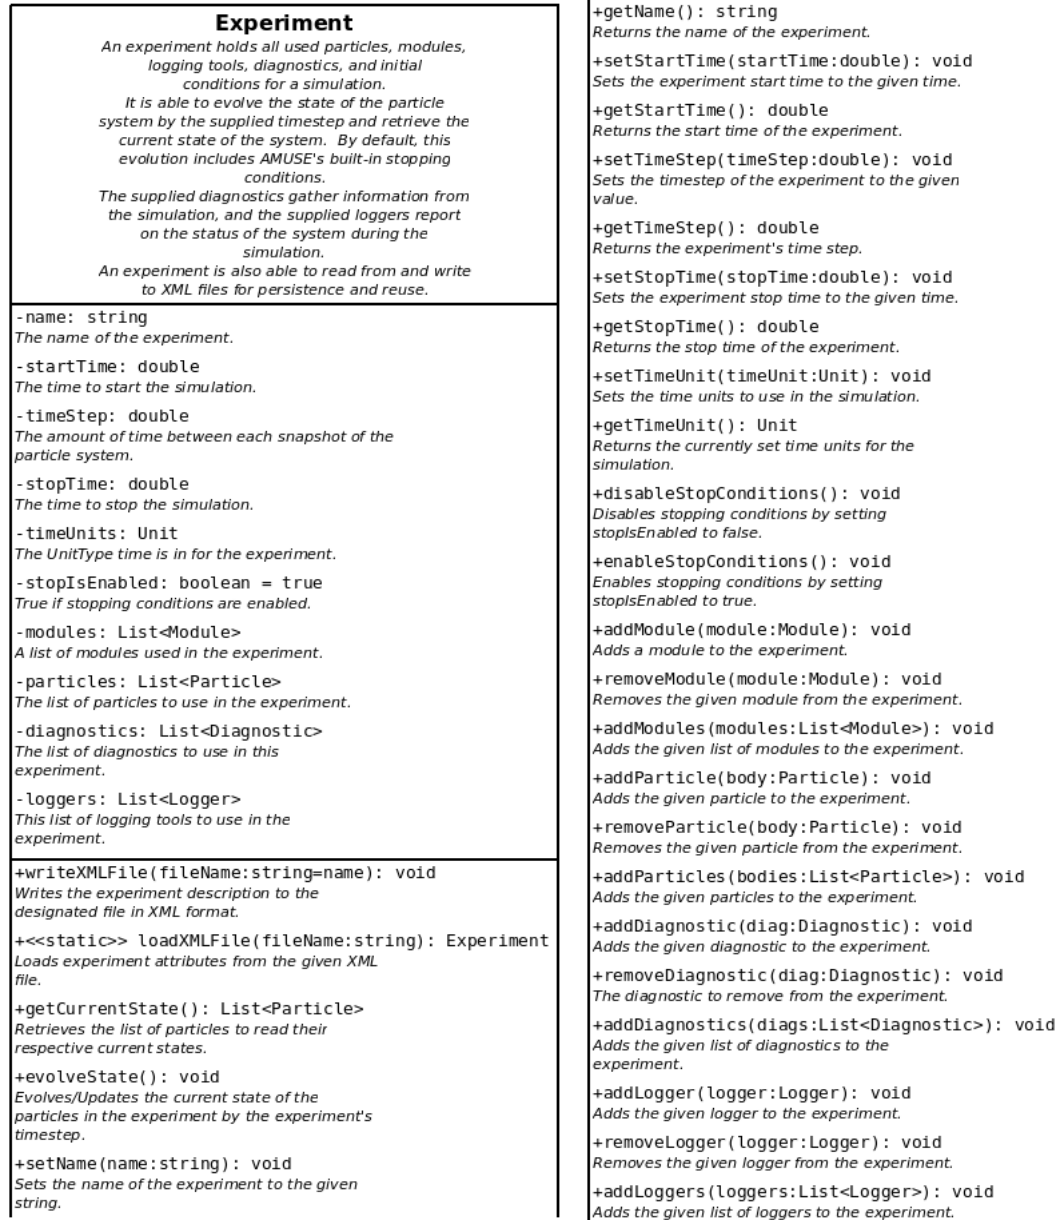


Figure 4.1: Experiment Class Diagram

4.1.1 Experiment Class

An instance of the Experiment class represents a single AMUSE experiment. It holds all particles (Satisfies requirement 2.3.3), modules, logging tools, diagnostics, and initial conditions for a simulation. These are described in more detail below. The experiment class can be viewed in Figure 4.1.

4.1.1.1 Experiment Attributes Experiments have the following attributes. The type UnitType is defined later in this document.

Name	Type	Description
name	string	The name of the experiment.
startTime	double	The start time of the experiment.
timeStep	double	The amount of time between each snapshot of the simulation.
stopTime	double	The stop time of the experiment.
timeUnit	Unit	The Unit (4.7) being used for time.
stopsEnabled	boolean	Whether or not stopping conditions are enabled.
modules	List<Module>	The list of Modules (4.4) used in the experiment.
particles	List<Particle>	The list of Particles (4.3) to use in the experiment.
diagnostics	List<Diagnostic>	The list of Diagnostics (4.13) to use in the experiment.
loggers	List<Logger>	The list of Loggers (4.14) to use in the experiment.

4.1.1.2 Experiment Operations

<i>void writeXMLFile(string fileName)</i>		<i>static Experiment loadXMLFile(string fileName)</i>	
Input	A string containing the filename to write the Experiment XML file to.	Input	A string containing the filename to load the Experiment XML from.
Output	(none)	Output	The Experiment specified by the XML file.
Description	Writes the Experiment in XML format to the specified filename.	Description	Recreates an Experiment from its XML specification.

<i>List<Particle></i> <i>getCurrentState()</i>		<i>void evolveState()</i>	
Input	(none)	Input	(none)
Output	The list of Particles in the Experiment.	Output	(none)
Description	Returns the list of Particles in the simulation with their updated attributes to read their current states.	Description	Evolves/Updates the current state of the particles in the Experiment by the Experiment's timeStep.
<i>void setName(string name)</i>		<i>string getName()</i>	
Input	A string containing the updated Experiment name.	Input	(none)
Output	(none)	Output	A string containing the name of the Experiment.
Description	Updates the Experiment's name to the given parameter.	Description	Returns the name of the Experiment.
<i>void setTimeUnit(Unit timeUnit)</i>		<i>Unit getTimeUnit()</i>	
Input	A Unit containing the updated Experiment timeUnit.	Input	(none)
Output	(none)	Output	A Unit containing the Experiment timeUnit.
Description	Updates the Experiment's timeUnit to the given parameter.	Description	Returns the timeUnit of the Experiment.
<i>void setStartTime(double startTime)</i>		<i>double getStartTime()</i>	
Input	A double containing the updated Experiment start time.	Input	(none)
Output	(none)	Output	A double containing the Experiment start time.
Description	Updates the Experiment's start time to the given parameter.	Description	Returns the start time of the Experiment.

<i>void setTimeStep(double timeStep)</i>		<i>double getTimeStep()</i>	
Input	A double containing the updated Experiment time step.	Input	(none)
Output	(none)	Output	A double containing the Experiment time step.
Description	Updates the Experiment's time step to the given parameter.	Description	Returns the time step of the Experiment.
<i>void setStopTime(double stopTime)</i>		<i>double getStopTime()</i>	
Input	A double containing the updated Experiment stop time.	Input	(none)
Output	(none)	Output	A double containing the Experiment stop time.
Description	Updates the Experiment's stop time to the given parameter.	Description	Returns the stop time of the Experiment.
<i>void disableStopConditions()</i>		<i>void enableStopConditions()</i>	
Input	(none)	Input	(none)
Output	(none)	Output	(none)
Description	Disables the Experiment's stopping conditions.	Description	Enables the Experiment's stopping conditions.
<i>void addModule(Module module)</i>		<i>void removeModule(Module module)</i>	
Input	A Module to add to the Experiment.	Input	A Module to remove from the Experiment.
Output	(none)	Output	(none)
Description	Adds the given Module to the Experiment.	Description	Removes the given Module from the Experiment.
<i>void addModules(List<Module> modules)</i>		<i>void addParticle(Particle body)</i>	
Input	A list of Modules to add to the Experiment.	Input	A Particle to add to the Experiment.
Output	(none)	Output	(none)
Description	Adds the given Modules to the Experiment.	Description	Adds the given Particle to the Experiment.

<i>void addParticles(List<Particle> bodies)</i>		<i>void removeParticle(Particle body)</i>	
Input	A list of Particles to add to the Experiment.	Input	A Particle to remove from the Experiment.
Output	(none)	Output	(none)
Description	Adds the given Particles to the Experiment.	Description	Removes the given Particle from the Experiment.
<i>void addDiagnostic(Diagnostic diag)</i>		<i>void addDiagnostics(List<Diagnostic> diags)</i>	
Input	A Diagnostic to add to the Experiment.	Input	A list of Diagnostics to add to the Experiment.
Output	(none)	Output	(none)
Description	Adds the given Diagnostic to the Experiment. Satisfies requirement 2.6.3.{1,2}.	Description	Adds the given Diagnostics to the Experiment.
<i>void removeDiagnostic(Diagnostic diag)</i>		<i>void addLogger(Logger logger)</i>	
Input	A Diagnostic to remove from the Experiment.	Input	A Logger to add to the Experiment.
Output	(none)	Output	(none)
Description	Removes the given Diagnostic from the Experiment.	Description	Adds the given Logger to the Experiment. Satisfies requirement 2.7.1.
<i>void addLoggers(List<Logger> loggers)</i>		<i>void removeLogger(Logger logger)</i>	
Input	A list of Loggers to add to the Experiment.	Input	A Logger to remove from the Experiment.
Output	(none)	Output	(none)
Description	Adds the given Loggers to the Experiment.	Description	Removes the given Logger from the Experiment. Satisfies requirement 2.7.2.

4.1.2 Experiment Hierarchy

An experiment hierarchy stores an experiment and all outputs of the runs of that experiment. Given this, the user may recreate an experiment or restart a failed experiment. The hierarchy may be stored in any format so long as there is an implementation of the methods in this interface to load and save the structure. Examples include databases, flat files, directory trees etc...

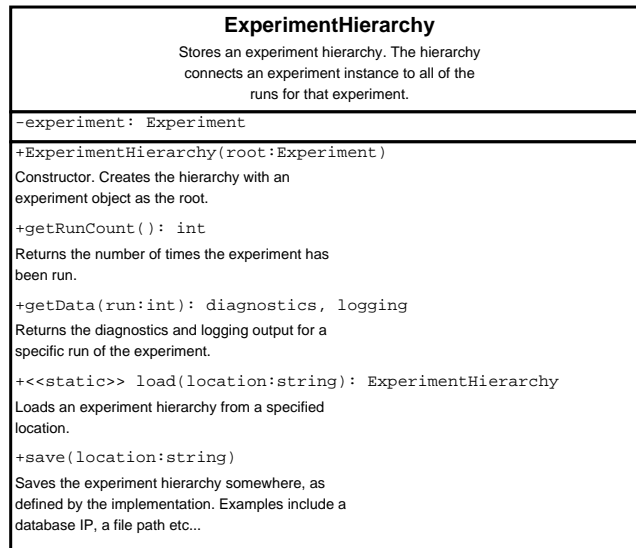


Figure 4.2: Experiment Hierarchy Class

4.2 Particles

Particles are the elements that are being updated and the experiment is being run on. These particles can be anything as little as atoms to as large as star clusters. They are made up of mass, position, and velocity. (Satisfies requirements 2.3.1, 2.3.2)

4.2.1 Particle Class




Figure 4.3: Particle Class Diagram

Particle class will be instantiated and used for storage of particle type items on the UI of GPUUnit. The Particle class in AMUSE will be used in the CLT. The particle class will be used primarily for data storage for experiment creation.

4.2.1.1 Particle Attributes

Name	Type	Description
mass	Parameter	The current mass of the particle. Will have units.
velocity	Parameter	The current mass of the particle. Will have units.
position	Parameter	The current position of the particle. Will have units.

4.2.1.2 Particle Methods

<i>Parameter getMass()</i>		<i>bool setMass(double mass, Unit unit)</i>	
Input	(none)	Input	The new value of the mass and the unit it is in.
Output	A parameter with the value and units of the current mass.	Output	Returns true if the mass is successfully set.
Description	Gets the current value of the mass.	Description	Sets the mass of the particle.
<i>Parameter getVelocity()</i>		<i>bool setVelocity(List<double> v, Unit unit)</i>	
Input	(none)	Input	The new vector value of the velocity and the unit it is in.
Output	A parameter with the value and units of the current velocity.	Output	Returns true if the velocity is successfully set.
Description	Gets the current value of the velocity.	Description	Sets the velocity of the particle.
<i>Parameter getPosition()</i>		<i>bool setPosition(List<double> pos, Unit unit)</i>	
Input	(none)	Input	The new vector value of the position and the unit it is in.
Output	A parameter with the value and units of the current position.	Output	Returns true if the position is successfully set.
Description	Gets the current value of the position.	Description	Sets the position of the particle.

<i>string toXml()</i>		<i>static Particle fromXml(string element)</i>	
Input	(none)	Input	A string containing an XML representation of a Particle.
Output	A string containing an XML representation of the Particle.	Output	A Particle whose property values are specified by the given XML element.
Description	Dumps the Particle to an XML element.	Description	Recreates a Particle from its XML specification.

4.3 Modules

This section defines an object model for the representation of and interaction with AMUSE modules, including several ancillary types utilized by the module implementation.

4.3.1 Module Class

Module
-name: string -description: string -domain: AstrophysicalDomain -codeName: string -codeLocation: string -isParallel: boolean -stoppingConditions: StoppingConditions -parameters: List<Parameter>
+getName(): string +setName(name:string): void +getDescription(): string +setDescription(description:string): void +getAstrophysicalDomain(): AstrophysicalDomain +setAstrophysicalDomain(domain:AstrophysicalDomain): void +getCodeName(): string +setCodeName(codeName:string): void +getCodeLocation(): string +setCodeLocation(codeLocation:string): void +getParallelism(): boolean +setParallelism(isParallel:boolean): void +getStoppingConditions(): StoppingConditions +setStoppingConditions(conds:StoppingConditions): void +getParameters(): List<Parameter> +setParameters(parameters:List<Parameter>): void +addParameter(p:Parameter): void +removeParameter(p:Parameter): boolean +toXml(): string +fromXml(element:string): Module

Figure 4.5: Module class diagram

An instance of the Module class represents a particular AMUSE module, providing an interface between GPUTask and the AMUSE code. The members of the Module class are used by GPUTask to properly locate and initialize an AMUSE module, as well as display a module's details and configuration to the user in the graphical interface. (Satisfies requirements 2.2.1, 2.2.2, 2.2.3, 2.2.4, 2.2.5)

4.3.1.1 Module Attributes

Name	Type	Description
name	string	The full name of the module.
description	string	A description of the module's purpose (in other words, the calculations performed by the module).
domain	AstrophysicalDomain	The astrophysical domain into which the module has been sorted.
codeName	string	The name of the AMUSE class containing the module code.
codeLocation	string	The location of the AMUSE module code.
isParallel	boolean	Whether the module's calculations can be parallelized across multiple workers by MPI.
stoppingConditions	StoppingConditions	The condition(s) under which the module may stop executing prematurely.
parameters	List<Parameter>	The module parameters. These module-specific values may be modified prior to running the experiment in order to fine-tune the module's behavior.

4.3.1.2 Module Operations

<i>string getName()</i>		<i>void setName(string name)</i>	
Input	(none)	Input	A new, full name for this module.
Output	The full name of the module.	Output	(none)
Description	Returns the name of the module.	Description	Updates this module's name to the given argument.

<i>string getDescription()</i>		<i>void setDescription(string description)</i>	
Input	(none)	Input	A new description of this module's purpose and calculations performed.
Output	A description of the module's purpose and the calculations it performs.	Output	(none)
Description	Returns this module's descriptive text.	Description	Updates this module's descriptive text to the given argument.
<i>AstrophysicalDomain getAstrophysicalDomain()</i>		<i>void setAstrophysicalDomain(AstrophysicalDomain domain)</i>	
Input	(none)	Input	An astrophysical domain into which this module will be sorted.
Output	The astrophysical domain into which the module has been sorted.	Output	(none)
Description	Returns this module's astrophysical domain.	Description	Updates this module's astrophysical domain.
<i>string getCodeName()</i>		<i>void setCodeName(string codeName)</i>	
Input	(none)	Input	The name of an AMUSE class containing this module's code.
Output	The name of the AMUSE class containing this module's code.	Output	(none)
Description	Returns the name of the module's associated AMUSE class.	Description	Updates this module's code reference to the given argument.
<i>string getCodeLocation()</i>		<i>void setCodeLocation(string codeLocation)</i>	
Input	(none)	Input	The location of the AMUSE code that defines this module.
Output	The location of the AMUSE code that specifies this module.	Output	(none)
Description	Returns the location of the module's AMUSE code.	Description	Updates this module's code location reference to the given argument.
<i>boolean getParallelism()</i>		<i>void setParallelism(boolean isParallel)</i>	
Input	(none)	Input	Whether the module's calculations can be parallelized across multiple workers by MPI.
Output	Whether the module's calculations can be parallelized across multiple workers by MPI.	Output	(none)
Description	Returns a boolean indicating whether the module's calculations can be parallelized.	Description	Sets or clears the parallelism flag for this module.

<i>StoppingConditions getStoppingConditions()</i>		<i>void setStoppingConditions(StoppingConditions conds)</i>	
Input	(none)	Input	A bitfield indicating the stopping conditions for this module.
Output	The conditions under which this module may stop executing prematurely.	Output	(none)
Description	Returns the stopping conditions specified for the module.	Description	Updates this module's stopping conditions to those represented by the given argument.
<i>List<Parameter> getParameters()</i>		<i>void setParameters(List<Parameter> parameters)</i>	
Input	(none)	Input	A list of module parameters.
Output	A list of the module's parameters.	Output	(none)
Description	Returns the parameters of this module.	Description	Updates this module's parameter list to the given argument.
<i>void addParameter(Parameter p)</i>		<i>boolean removeParameter(Parameter p)</i>	
Input	A module parameter to add to the module.	Input	A module parameter to remove from the module.
Output	(none)	Output	Whether the given parameter was removed. Returns false if the parameter was not found in this module's parameter list.
Description	Adds the given parameter to this module.	Description	Removes the given parameter from this module, if the parameter exists.
<i>string toXml()</i>		<i>static Module fromXml(string element)</i>	
Input	(none)	Input	A string containing an XML representation of a Module.
Output	A string containing an XML representation of the Module.	Output	A Module whose property values are specified by the given XML element.
Description	Dumps the Module to an XML element.	Description	Recreates a Module from its XML specification.

4.3.2 Parameter Class

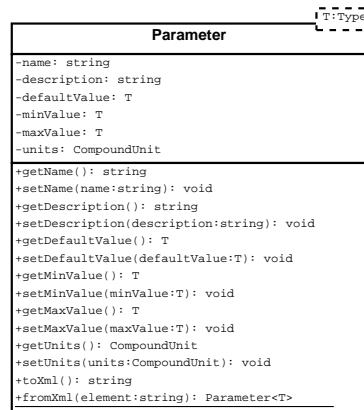


Figure 4.6: Parameter class diagram

Modules have parameters that may be modified by the user via the graphical interface. These parameters are typically specific to the module's domain or even to the individual module. The Parameter class is a generic class whose type parameter specifies the type of the parameter's value (e.g. integer or floating-point). A parameter may be a physical quantity, flag, or other value. (Satisfies Requirement 2.3.6)

4.3.2.1 Parameter Attributes

Name	Type	Description
T	Type	The data type of the parameter's value.
name	string	The brief, descriptive name of the parameter.
description	string	A description of the parameter's meaning and effects.
defaultValue	T	The parameter's default value.
minValue	T	A lower bound on the range of valid parameter values.
maxValue	T	An upper bound on the range of valid parameter values.
units	CompoundUnit	The physical unit(s) associated with the parameter's value.

4.3.2.2 Parameter Operations

<i>string getName()</i>		<i>void setName(string name)</i>	
Input	(none)	Input	A new name for this parameter.
Output	The brief, descriptive name of the parameter.	Output	(none)
Description	Returns the name of the parameter.	Description	Updates this parameter's name to the given argument.

<i>string getDescription()</i>		<i>void setDescription(string description)</i>	
Input	(none)	Input	A new description for this parameter.
Output	A description of the parameter's meaning and effects.	Output	(none)
Description	Returns this parameter's descriptive text.	Description	Updates this parameter's descriptive text to the given argument.
<i>T getDefaultValue()</i>		<i>void setDefaultValue(T defaultValue)</i>	
Input	(none)	Input	A new default value for this parameter.
Output	The parameter's default value.	Output	(none)
Description	Returns the default value for this parameter.	Description	Updates this parameter's default value to the given argument.
<i>T getMinValue()</i>		<i>void setMinValue(T minValue)</i>	
Input	(none)	Input	A new lower bound on this parameter's range of possible values.
Output	The parameter's minimum valid value.	Output	(none)
Description	Returns the minimal allowed value for this parameter.	Description	Updates this parameter's minimum value to the given argument.
<i>T getMaxValue()</i>		<i>void setMaxValue(T maxValue)</i>	
Input	(none)	Input	A new upper bound on this parameter's range of possible values.
Output	The parameter's maximum valid value.	Output	(none)
Description	Returns the maximal allowed value for this parameter.	Description	Updates this parameter's maximum value to the given argument.
<i>CompoundUnit getUnits()</i>		<i>void setUnits(CompoundUnit units)</i>	
Input	(none)	Input	A new compound physical unit to be associated with this parameter.
Output	An object representing the physical unit(s) associated with the parameter's value.	Output	(none)
Description	Returns an object containing the units associated with this parameter.	Description	Updates this parameter's units to the given argument.

<i>string toXml()</i>		<i>static Parameter<T> fromXml(string element)</i>	
Input	(none)	Input	A string containing an XML representation of a Parameter.
Output	A string containing an XML representation of the Parameter.	Output	A Parameter whose property values are specified by the given XML element.
Description	Dumps the Parameter to an XML element.	Description	Recreates a Parameter from its XML specification.

4.3.3 Unit Class




Figure 4.7: Unit class diagram

With AMUSE, physical quantities may be expressed using, and converted between, a number of different standard units. The Unit class, in conjunction with the CompoundUnit class, allows GPUUnit to present these units to the user for the purposes of description and selection.

4.3.3.1 Unit Attributes

Name	Type	Description
type	UnitType	The base type of astrophysical unit being represented.
prefix	SIPrefix	The SI prefix that modifies the base unit's order of magnitude.
exponent	integer	The exponent to be applied to the unit. For example, a typical unit of volume (such as cubic meters) would have an exponent of 3.

4.3.3.2 Unit Operations

<i>UnitType</i> <i>getType()</i>		<i>void setType(UnitType type)</i>	
Input	(none)	Input	A base astrophysical unit to which this instance should be updated.
Output	The base type of astrophysical unit being represented.	Output	(none)
Description	Returns the base type of astrophysical unit represented by this instance.	Description	Updates this instance's base unit type to the given argument.
<i>SIPrefix</i> <i>getPrefix()</i>		<i>void setPrefix(SIPrefix prefix)</i>	
Input	(none)	Input	An SI prefix that will replace this instance's current SI prefix.
Output	The SI prefix that modifies the base unit's order of magnitude.	Output	(none)
Description	Returns the SI prefix that augments the magnitude of this unit.	Description	Updates this instance's SI prefix to the given argument.
<i>UnitType</i> <i>getExponent()</i>		<i>void setExponent(UnitType type)</i>	
Input	(none)	Input	A new exponent to be applied to this unit.
Output	The exponent applied to this unit.	Output	(none)
Description	Returns the exponent applied to this unit.	Description	Updates this instance's exponent to the given argument.

4.3.4 CompoundUnit Class

CompoundUnit
<pre> ~description: string ~symbolicDescription: string ~units: List<Unit> +getDescription(): string +setDescription(description:string): void +getSymbolicDescription(): string +setSymbolicDescription(symbolicDescription:string): void +getUnits(): List<Unit> +setUnits(units:List<Unit>): void +addUnit(u:Unit): void +removeUnit(u:Unit): boolean </pre>

Figure 4.8: CompoundUnit class diagram

When performing physical calculations, units of measure may be combined in any number of ways. The CompoundUnit class provides the mechanism by which units may be presented to the user, regardless of whether they are simple or compound.

4.3.4.1 CompoundUnit Attributes

Name	Type	Description
description	string	A non-abbreviated textual description of the combined physical unit.
symbolicDescription	string	A shorthand textual description of the combined physical unit, using unit abbreviations.
units	List<Unit>	The list of simple units whose combination is represented by this compound unit.

4.3.4.2 CompoundUnit Operations

<i>string getDescription()</i>		<i>void setDescription(string description)</i>	
Input	(none)	Input	A new, full description for this compound unit.
Output	A non-abbreviated textual description of the combined physical unit.	Output	(none)
Description	Returns a textual description of this combined unit.	Description	Updates this instance's description to the given argument.
<i>string getSymbolicDescription()</i>		<i>void setSymbolicDescription(string symbolicDescription)</i>	
Input	(none)	Input	A new abbreviated description for this compound unit.
Output	A shorthand textual description of the combined physical unit, using unit abbreviations.	Output	(none)
Description	Returns an abbreviated description of this combined unit.	Description	Updates this instance's short description to the given argument.
<i>List<Unit> getUnits()</i>		<i>void setUnits(List<Unit> units)</i>	
Input	(none)	Input	A list of simple units to which this instance should be updated.
Output	The list of simple units whose combination is represented by this compound unit.	Output	(none)
Description	Returns the list of units contained in this compound unit.	Description	Updates this instance's list of simple units to the given argument.

<i>void addUnit(Unit u)</i>		<i>boolean removeUnit(Unit u)</i>	
Input	A unit to be included in this compound unit.	Input	A unit to remove from this compound unit.
Output	(none)	Output	Whether the given unit was removed. Returns false if the unit was not found in this instance's unit list.
Description	Adds the given unit to this compound unit's list of simple units.	Description	Removes the given simple unit from this compound unit, if the simple unit exists.

4.3.5 UnitType Enumeration




Figure 4.9: UnitType type diagram

UnitType enumerates the myriad of base physical units supported by AMUSE. The list of members shown may not be exhaustive.

4.3.6 SIPrefix Enumeration




Figure 4.10: SIPrefix type diagram

When dealing with physical quantities, standard prefixes may be prepended to units to denote a given quantity's order of magnitude. SIPrefix enumerates the possible unit prefixes.

4.3.7 AstrophysicalDomain Enumeration




Figure 4.11: AstrophysicalDomain type diagram

The modules included with AMUSE have categorized according to their domain. A module's domain indicates the quantities it deals with as well as common operations that may be called on the module. The AstrophysicalDomain type enumerates the several domains that have been specified by AMUSE.

4.3.8 StoppingConditions Enumeration

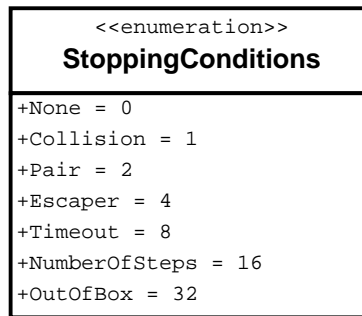


Figure 4.12: StoppingConditions type diagram

StoppingConditions is a set of bit flags enumerating the conditions under which a module may return from execution before completing its assigned calculations.

4.4 Diagnostics

Diagnostics have access to the state of the experiment at the end of each time step. Default diagnostics will be provided, e.g. printing the state to a file at a specified time step or creating a stellar mass histogram. Users can create custom diagnostic plug-ins. (Satisfies requirement 2.6)

4.4.1 Object Model

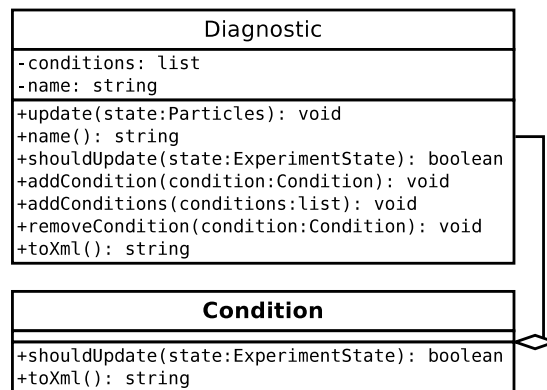


Figure 4.13: Diagnostics Object Model

4.4.2 Diagnostic Class

The Diagnostic class is an abstract class that is the base of both built-in and custom diagnostics tools. All custom diagnostic scripts must inherit from this class. Subclasses must implement the update and name operations. The Diagnostic class also contains a list of conditions to check if the diagnostic needs to be updated.

4.4.3 Experiment Manager Class

The Experiment Manager class has the responsibility to manage as well as update the active diagnostics.

4.4.4 Condition Class

The Condition class represents the conditions that need to be satisfied for a diagnostic to be updated.

4.5 Logging

Logs are helpful when the experiment must be restarted along with its output data if the experiment has an error or needs to be stopped for any other reason. Logging could be enabled or disabled by the user. The output destination for the logs could be specified. The specified destination for the output can be one among the following: the console, a file, or a window in the GUI. (Satisfies requirement 2.7)

4.5.1 Object Model




Figure 4.14: Logger Object Model

4.5.2 Logger Class

The Logger class takes care of selecting the required output option as the console, file or the GUI window.

5 Experiment Generation

5.1 Overview

Part of the framework includes generating the Python scripts that will run the experiment. Experiment generation by default will produce a script as defined in Section 6. The user may customize this generation process by providing custom code that is run prior to module initialization, before, and after each timestep. The user can also create custom Logging and Diagnostic modules which will be run by the generated experiment.

ExperimentGenerator This class is responsible for generating Python code to run a specific experiment. This includes module initialization, particle setup etc...
<code>-customCodes: map<Module, string></code> Contains optional custom code to initialize a module. If no custom code is given, modules are initialized based on the predefined rules described in section 5 for the command line tool.
<code>+generate(experiment:Experiment): string</code> Returns a string containing the complete Python code to run the provided experiment. <code>+addCustomModuleCode(code:string,module:Module)</code> Adds custom code to be called before initializing a module. <code>+addPreEvolutionCode(code:string)</code> Adds custom python code to be run prior to evolving the model. <code>+addPostEvolutionCode(code:string)</code> Adds custom python code to be run after evolving the model.

Figure 5.1: Experiment Generation

6 Command Line Tool

6.1 Overview

The command line tool (CLT) is used to run an experiment file. The tool is the end result of the code generation process. The CLT begins by parsing the experiment file. After parsing the file, the CLT takes the initialization parameters and loads them using the AMUSE framework. The CLT then takes the modules specified, imports them from the available AMUSE modules and initializes them. The CLT then links the modules together to allow interaction and updating of the current state of the system using AMUSE channels. Next the CLT loads the Logging and Diagnostic scripts specified by the experiment file, linking it to the particle list.

The CLT then runs through the timesteps specified in the experiment file, evolving all of the modules loaded at each time step. It will run the diagnostic scripts as well as logging on the intervals specified. After the model is done evolving it will output using whatever diagnostic scripts listed, in the format specified by the script.

6.2 Using the CLT

The command line tool can be used by either a user or the GUI to send out jobs and run them. The command line tool though can be used with a basic experiment file with some on the fly modifications by using some flags. (Satisfies requirement 2.4.1)

6.2.1 CLT Flags

Flag	<Value Being Passed>	Description
-help	None	Displays a useful help prompt with the list of CLT flags.
-f	Experiment File Location	Loads the experiment file prompted by the filename.
-n	Number of Particles	Changes the number of particles being passed in the experiment file.
-t	End Time	Changes the max time allowed in the simulation.
-dt	Timestep	Changes the timestep in which dt moves forward.
-r	Radius	Changes the radius scaling factor.

6.3 Running The Experiment

6.3.1 Simulation Overview

A Simulation is the goal of both GPUUnit and AMUSE, which makes it important to understand and break it into key parts. The Simulation can be looked at as three sections: Simulation Initialization, Running Simulation, Post Simulation. This section goes into more detail about all three of these parts and what happens in each. Each section title that corresponds with a section in figure 6.1 will have a hyphenated number to correspond with a step in the flow chart.




Figure 6.1: Flow Chart of Typical AMUSE Simulation

6.3.1.1 Simulation Initialization Simulation Initialization is used for preparation of the particles and modules to be able to model the system. The following breaks down the initialization into the main steps.

6.3.1.1.1 Instantiate Particles - 1 In the initialization phase of the experiment it is important to select proper initial conditions. The main initial conditions to set are the mass distribution and starting locations of the particles. There are built in AMUSE classes that are used to set these. The CLT will find the initial condition files by name while reading in the experiment file. (Satisfies Requirements 2.3.4)

6.3.1.1.2 Create Conversion Object - 2 Most astronomical simulations use a scaling factor to hasten numerical integration. A conversion object is a class in AMUSE instantiated by passing in the total mass and average radius of the system. This conversion object is passed into all modules each module knows how to perform unit conversions from unitless numbers.

6.3.1.1.3 Module Creation - 3.4 The CLT instantiates the AMUSE class for the specific module it is working on from the location specified in the module class. The module is instantiated with a conversion object being passed into its constructor. This conversion object is used for the aforementioned tasks in section 6.3.1.1.2. The module then sets its list of particles to the master list of particles. After, it runs its initial calculations needed for module evolution.

6.3.1.1.4 Instantiate Channels - 5 Channels are used for communication of the data to and from the modules that may be on a different system. At each time step, the master list of particles is synchronized to each module's separate list of particles. After, the modules' particles are re-synchronized to the master list of particles. This allows the modules to have the same data as the simulation moves forward. (Satisfies Requirement 2.5.1, 2.5.2, 2.5.2.1)

6.3.1.2 Running Simulation Running the system is the area inside of the while loop. This is where the modules are run, logging and diagnostics are performed, and where the time scale increments. (Satisfies Requirement 2.4.2)

6.3.1.2.1 Evolve Modules - 6 The evolution of the Module is the core of running a simulation. The evolution of the module is done by running the AMUSE module's *evolve_model* function with the current time.

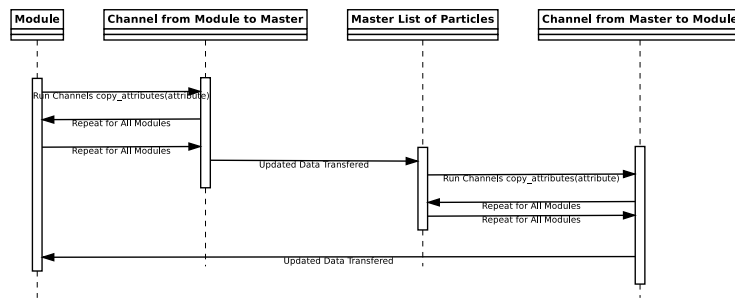


Figure 6.2: Synchronizatn Phase

6.3.1.2.2 Synchronize Particles - 7 Particle synchronization is required because every module has a local list of particles to allow for parallelization. Since at most only one of each module type will be run per experiment, there will be no

overlapping data modifications. After data is updated, we use the aforementioned channels to do the synchronization, first from the master list of particles, then back (As shown in figure 6.2). This ensures all modules have the latest version of particle data.

6.3.1.2.3 Run Diagnostics - 8 The function loops through all diagnostic scripts, running their *shouldUpdate* method with the current state to see if they should run. If a diagnostic should run, it then runs the *update* method passing in the master list of particles.(Figure 4.13)

6.3.1.2.4 Run Logging - 9 The function loops through all logging scripts and runs any real time logging scripts that need to be run. The real time logging scripts are denoted by name and run their *logData* method passing in the list of particles.(Figure 4.14)

6.3.1.2.5 Increment Time - 10 The time increments using the specified timestep (Figure 4.1). The experiment stops when the stopping conditions are met. (Satisfies Requirement 2.4.3.1)

6.3.1.3 Post Simulation

6.3.1.3.1 Run Closing Diagnostics Scripts - 11 The function loops through all diagnostic scripts, running their *shouldUpdate* method with the current state to see if they should run. If the diagnostic should run, it then runs the update method passing in the master list of particles.(Figure 4.13)

6.3.1.3.2 Run Closing Logging Scripts - 12 The function loops through all logging scripts and runs any real time logging scripts that need to be run. The real time logging scripts are denoted by name and run their *logData* method passing in the list of particles.(Figure 4.14)

6.3.1.3.3 Stop Modules - 13 The modules are cleaned up after being run. If running in parallel, a kill command is sent via MPI. We will run the modules'(AMUSE class) *stop()* command.

7 Networking

7.1 Overview

The networking layer is responsible for communication between the interface and the cluster status daemons running on the cluster nodes. The user interface and console can send messages through the network which will be distributed via IP multicast to any cluster nodes listening to a selected multicast group. Cluster nodes will be able to send replies to the sender containing information such as health stats and the

number of experiments running on that node. The UI and command line tools will use IPC (pipes etc...) to communicate with local instances of the control program to send network packets to the cluster nodes.

7.2 Network Software Components

7.2.1 Control Instance

A control instance is a program that will run in the background on any machine wishing to control experiments on a cluster. The core scripts (command line tools) and UI will both communicate with this instance to perform any network access.

7.2.2 Node Instance

Node instances are processes that run on each computing node in the cluster. The node instance tracks health statistics and running experiments and can send these pieces of information as direct (unicast) replies to queries received through the multicast group.

7.3 Object Model

7.3.1 Overview

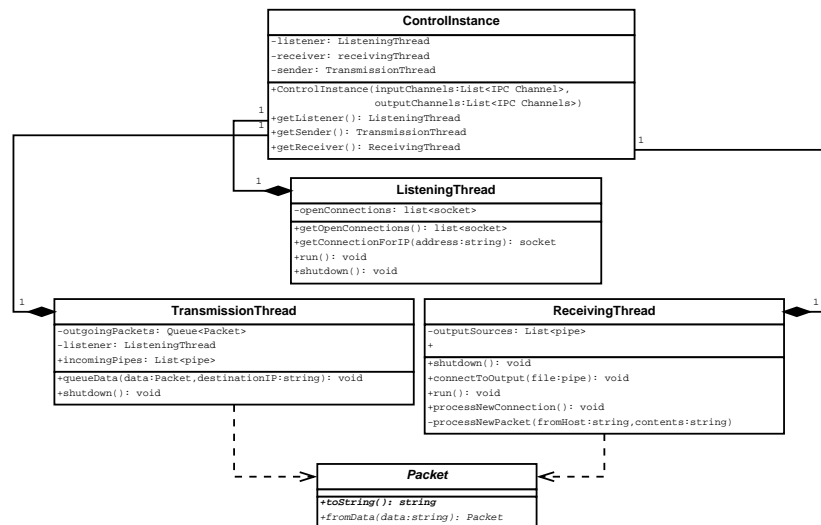


Figure 7.1: Network Object Model

7.3.2 Control Instance

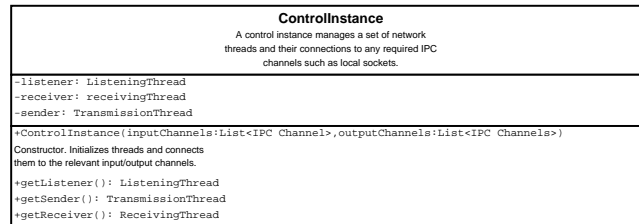


Figure 7.2: Control Instance Class Diagram

7.3.3 Node Instance

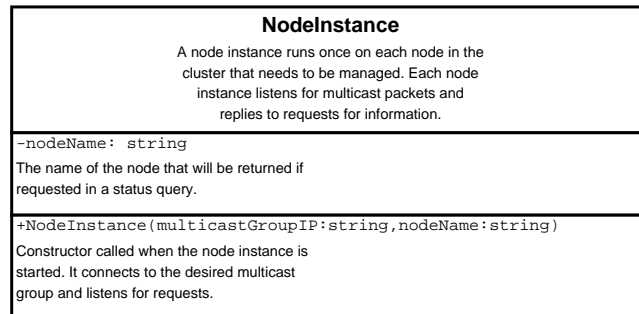


Figure 7.3: Node Instance Class Diagram

7.3.4 TransmissionThread

The transmission thread (Figure 7.4) will process outgoing packets queued for transmission by external sources such as the UI.

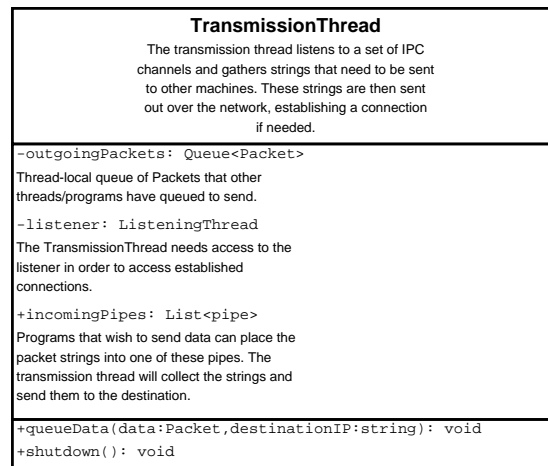


Figure 7.4: Transmission Thread Class Diagram

7.3.5 ReceivingThread

The receiving thread (Figure 7.5) will read incoming packets from established connections with node instances.

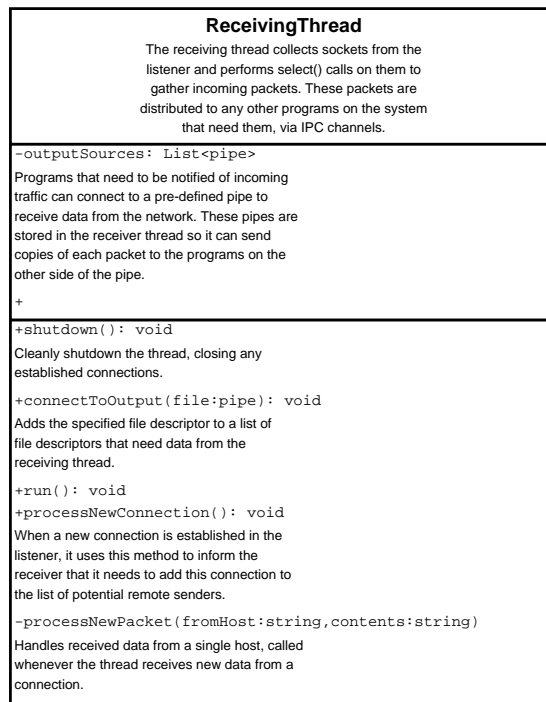


Figure 7.5: Receiving Thread Class Diagram

7.3.6 ListeningThread

The listening thread (Figure 7.6) will listen for and accept new incoming connections from node instances.

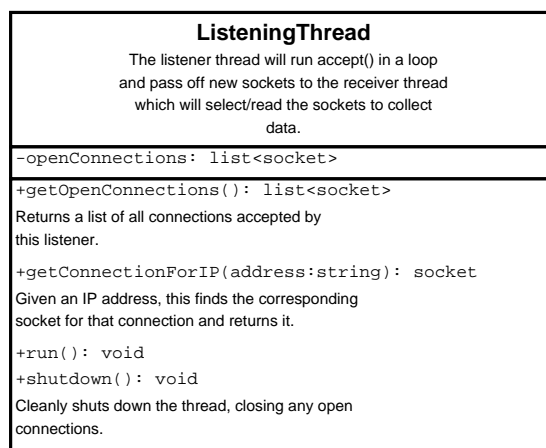


Figure 7.6: Listening Thread Class Diagram

7.3.7 Packet Objects

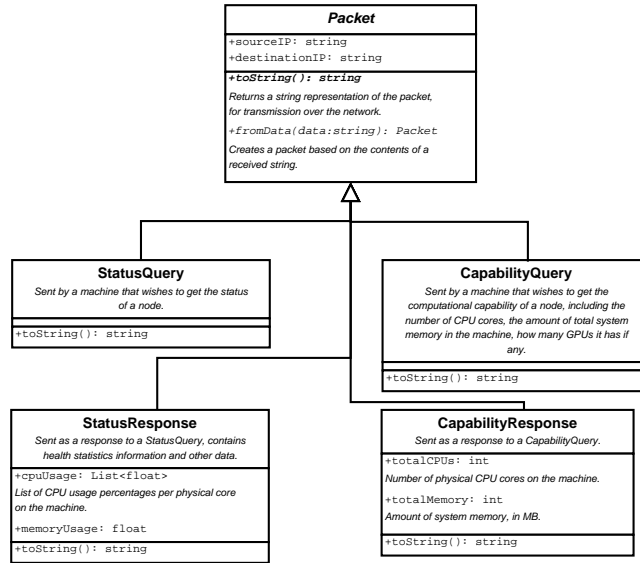


Figure 7.7: Packet Classes

8 Data

8.1 Experiment Specification

The following is a list of XML specifications for experiments. Experiments are saved and loaded via this format. Note that not all specified attributes will always be present because not all attributes are relevant to every experiment. See Figure 8.1 for an example XML experiment file. (Satisfies requirement 2.3.3, 2.3.5)

```

<experiment name="myExperiment" stopEnabled="true">
  <time units="Myr" start="0" step="300" end="900"/>
  <module type="BHtree">
    <param name="use_self_gravity" value="0"/>
    <stopCondition type="COLLISION_DETECTION"/>
  </module>
  <particle radius="6371|km" mass="5.9736e4|kg" position="[1,0,0]|AU"
velocity="[0.0,29783,0.0]|ms"/>
  <particle radius="1|RSun" mass="1|MSun" position="[0,0,0]|m"
velocity="[0.0,0.0,0.0]|ms"/>
  <diagnostic name="massStats" file="massDiag.py"/>
  <diagnostic name="velocityStats" file="vDiag.py"/>
  <logger name="stateUpdates" file="stateUpdates.py"/>
  <logger name="timeSteps" file="timeSteps.py"/>
</experiment>
  
```

Figure 8.1: Example XML Experiment File

8.1.1 Experiment

Experiment tags are the first and last tags in an experiment specification because they encapsulate all experiment attributes.

Attribute	Description
name	The name of the experiment
stopEnabled	Whether or not stopping conditions are enabled: “true” or “false”

8.1.2 Time

Time tags specify the time parameters of the experiment.

Attribute	Description
units	The time units for the experiment to use, defined by AMUSE
start	The start time of the experiment
step	The amount of time between simulation snapshots
end	The time to end the experiment

8.1.3 Module

Module tags define the modules used in the experiment. Each module has child param tags to define initial conditions. These initial conditions are defined by AMUSE.

Attribute	Description
name	The name of the module to use

8.1.4 Param

Param tags define the parameters for their parent model. Each defines a model's parameter value for use in the experiment. All parameters, default values, units, and descriptions are defined by AMUSE.

Attribute	Description
name	The name of the parameter to initialize
value	The value to give the parameter

8.1.5 Particle

Particle tags define the particles used in the experiment. All references to units are those defined by AMUSE.

Attribute	Description
mass	The mass of the particle with units
radius	The radius of the particle with units
position	The position of the particle in [X,Y,Z] format with units
velocity	The velocity of the particle in [X,Y,Z] format with units
luminosity	The luminosity of the particle with units
temperature	The temperature of the particle with units
age	The age of the particle with units
stellarType	The stellar type of the particle, defined by AMUSE

8.1.6 Diagnostic

Diagnostic tags define the diagnostic scripts used in the experiment.

Attribute	Description
name	The name of the diagnostic for reference purposes
file	The filename of the diagnostic script

8.1.7 Logger

Logger tags define the logging scripts used in the experiment.

Attribute	Description
name	The name of the logger for reference purposes
file	The filename of the logging script

8.1.8 Stopping Condition

Stopping Condition tags define the stopping conditions used in the experiment. All stopping conditions are defined by AMUSE.(Satisfies Requirement 2.4.3, 2.4.3.2)

Attribute	Description
type	The type of stopping condition to use

8.2 Experiment Results Directory Structure

The following is a list of specifications for GPUUnit's default directory structure for storing experiment results. Notice that it is three-tiered for ease in comparing results from multiple runs of the same experiment.

8.2.1 Top Level Directory

By default, GPUUnit creates a directory named "Experiments" in its current location. The user is also able to change this location and/or choose a different directory for storing experiment data. This directory contains a directory for each experiment created and saved by GPUUnit. These directory names match the name of each experiment, dictated by the user during a save operation.

8.2.2 Single Experiment Directory

Each experiment directory contains a directory corresponding to a previously run or currently running simulation. By default, these are named "Run #" where '#' is a number corresponding to each run of the experiment. These values start at 1 and continue ad infinitum. e.g. "Run 1, Run 2, Run 3, ..." where the commas separate directory names.

8.2.3 Run Directory

Each "Run #" directory contains the initial XML specification for the experiment and the results from logging and diagnostics.

8.3 Network Packets

The following is a list of packet data content specifications. Data is described in terms of types to ease implementation, however the information is sent over the network as a string to avoid any machine-specific byte ordering issues. Fields in the packet are separated by a "|" character. If this character is to appear as text inside a packet for any reason, it must be escaped as "\\|"

8.3.1 Packet Header

The packet header precedes data in all packets.

PACKET_TYPE	LENGTH	SOURCE_IP	DEST_IP
int	long	string	string

8.3.2 Status Query Packet

The status request may contain flags requesting additional data from the node beyond the elements specified here. If the recipient understands the flags, they will fill the ADDITIONAL_DATA field with the appropriate response.

HEADER	ADDITIONAL_REQUEST_FLAGS
-	string

8.3.3 Status Response Packet

HEADER	CPU_USAGE	MEMORY_USAGE	SIMS_RUNNING	ADDITIONAL_DATA
-	float	float	string	string

8.3.4 Capability Query Packet

The additional data protocol is the same here as in the status request above.

HEADER	ADDITIONAL_REQUEST_FLAGS
-	string

8.3.5 Capability Response Packet

HEADER	NUM_CPUS	MAX_MEMORY	NUM_GPUS	ADDITIONAL_DATA
-	int	long	int	string

9 Appendix

9.1 Requirements Traceability Table

Requirement	Section in SDD
2.2.1	4.3.1
2.2.2	4.3.1
2.2.3	4.3.1
2.2.4	4.3.1
2.2.5	4.3.1
2.3.1	4.2
2.3.2	4.2
2.3.3	4.1
2.3.4	6.3.1.1.1
2.3.5	7.1
2.3.6	4.3.2
2.4.1	6.2
2.4.2	6.3.1.2
2.4.3	6.3.1.2.5
2.4.4	6.3.1.2.5
2.5.1	6.3.1.1.4
2.5.2	6.3.1.1.4
2.6.1	4.4
2.6.2	4.4
2.6.3	4.4
2.6.4	4.4
2.6.5	4.4
2.6.6	4.4
2.7.1	4.5
2.7.2	4.5
2.7.3	4.5
2.8.1	3.3
2.8.2	3.3

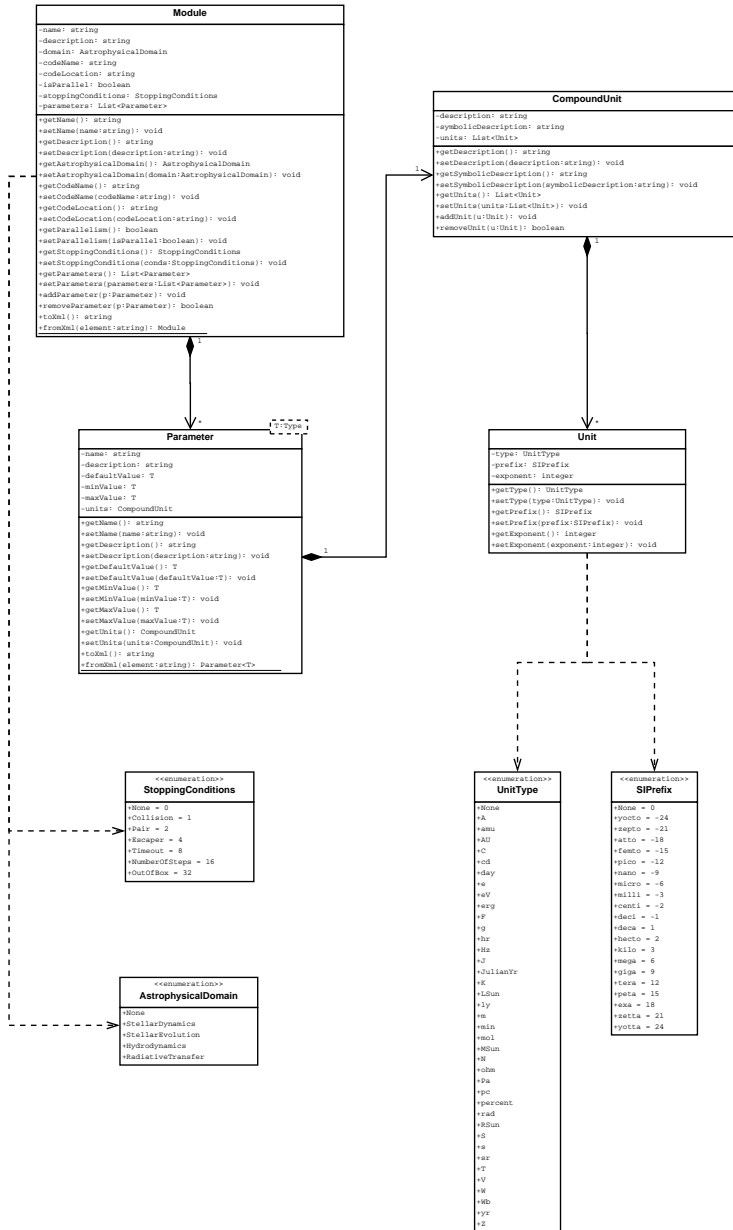


Figure 4.4: Relationships between Module and supporting classes