

Software Design Document:

GPU N-Body Integration Toolkit

Prepared by:

Daniel Bagnell

Jason Economou

Rajkumar Jayachandran

Tim McJilton

Gabe Schwartz

Andrew Sherman

Advisor:

Prof. Jeremy Johnson

Stakeholders:

Prof. Steve McMillan

Alfred Whitehead

Contents

1	Introduction	4
1.1	Purpose	4
1.2	Scope	4
1.3	Glossary	4
2	Architecture	5
3	Graphical User Interface	5
3.1	Overview	5
3.2	Network Interaction	5
3.3	Node Management	6
4	Experiment Components	8
4.1	Experiment Object Model	8
4.1.1	Experiment Class	8
4.2	Particles	11
4.2.1	Particle Class	12
4.3	Modules	13
4.3.1	Module Class	13
4.3.2	Parameter Class	15
4.3.3	Unit Class	16
4.3.4	CompoundUnit Class	16
4.3.5	UnitType Enumeration	18
4.3.6	SIPrefix Enumeration	19
4.3.7	AstrophysicalDomain Enumeration	20
4.3.8	StoppingConditions Enumeration	20
4.4	Diagnostics	20
4.4.1	Object Model	21
4.4.2	Diagnostic Class	21
4.4.3	Experiment Manager Class	21
4.4.4	Condition Class	21
4.5	Logging	21
4.5.1	Object Model	22
4.5.2	Logging Class	22
5	Command Line Tool	22
5.1	Overview	22
5.2	Using the CLT	22
5.2.1	CLT Flags	23
5.3	Running The Experiment	23
5.3.1	Simulation Overview	23

6	Networking	26
6.1	Overview	26
6.2	Network Software Components	26
6.2.1	Control Instance	26
6.2.2	Node Instance	27
6.3	Object Model	27
6.3.1	Overview	27
6.3.2	Control Instance	27
6.3.3	Node Instance	28
6.3.4	TransmissionThread	28
6.3.5	ReceivingThread	28
6.3.6	ListeningThread	29
6.3.7	Packet Objects	30
7	Data	30
7.1	Experiment Specification	30
7.1.1	Experiment	31
7.1.2	Time	31
7.1.3	Module	31
7.1.4	Param	31
7.1.5	Particle	31
7.1.6	Diagnostic	32
7.1.7	Logger	32
7.1.8	Stopping Condition	32
7.2	Network Packets	32
7.2.1	Packet Header	32
7.2.2	Status Query Packet	32
7.2.3	Status Response Packet	32
7.2.4	Capability Query Packet	33
7.2.5	Capability Response Packet	33
8	Appendix	33
8.1	Requirements Traceability Table	33

1 Introduction

1.1 Purpose

This document serves as the system design for the GUnit framework to control AMUSE experiments as set forth in the Software Requirements Specification. AMUSE (Astrophysical Multipurpose Software Environment) is a software package and set of legacy codes for performing astrophysics simulations. For more details on the purpose and use cases of AMUSE, please refer to our SRS. The goal of this document is to provide details guiding the construction of the framework, as well as to provide a reference to the framework's architecture. It contains specifications for the objects used by the system and our initial reference implementation plans.

1.2 Scope

This document will describe the implementation for our framework, covering the network layer, the user interface and the core python command line scripts. It will cover user interface support classes as they will need to be coded, however the actual UI components such as windows and widgets will not be described. Code and objects for these items are generated as part of the Qt Creator¹ GUI development kit. As such, the code is not intended to be human-readable and is subject to change between Qt versions. Our SRS document contains the reference prototype which will serve as the design for the GUI widgets. The data models including experiment file layouts and network packet details will also be covered.

1.3 Glossary

For terms not in this glossary, please refer to the SRS.

Cluster A group of networked computing nodes available to perform some task.

Computer Node A computer that is a member of a cluster.

IPC Channel Any object or memory space used for Inter-Process Communication (IPC). Examples include UNIX Pipes, BSD/Winsock Sockets and UNIX Shared Memory (SHM).

Node Health Statistics Information about the current state of a node including properties such as CPU usage and memory usage.

¹<http://qt.nokia.com/products/developer-tools/>

2 Architecture

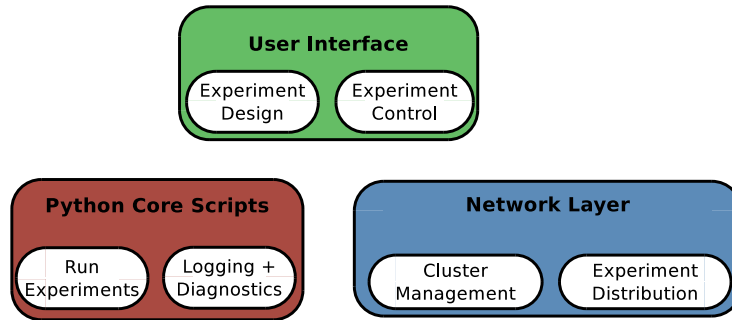


Figure 2.1: Architecture Diagram

3 Graphical User Interface

3.1 Overview

The user interface will allow the user to create and manage experiments as described in the SRS. Much of the interface code is generated from the GUI development tools, however there are some support classes required for the GUI to interact with the network and the command line tools.

3.2 Network Interaction

The network interaction classes allow the UI to send and receive messages through a control instance (Section 6.2.1). The control instance can be either running locally or on a remote machine. The UI uses this link to query the cluster for status updates and control running experiments.

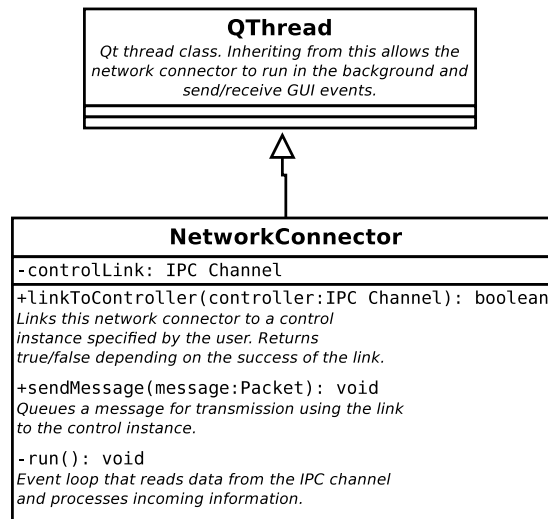


Figure 3.1: GUI Network Connection Class

3.3 Node Management

Nodes in the cluster, detected through the network connector, may be viewed in the node view window. (Satisfies requirements 2.8.1.X)

From the node window, the user can assign experiments to individual nodes. (Satisfies requirement 2.8.2.1)

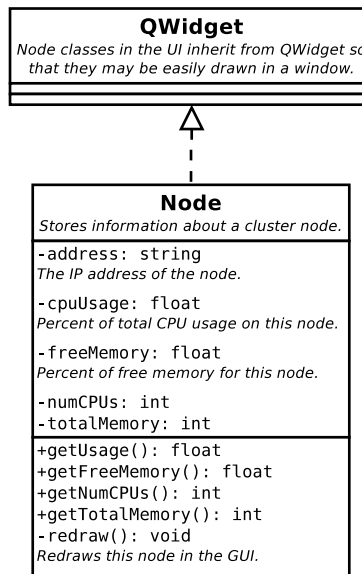


Figure 3.2: Node Class

4 Experiment Components

4.1 Experiment Object Model

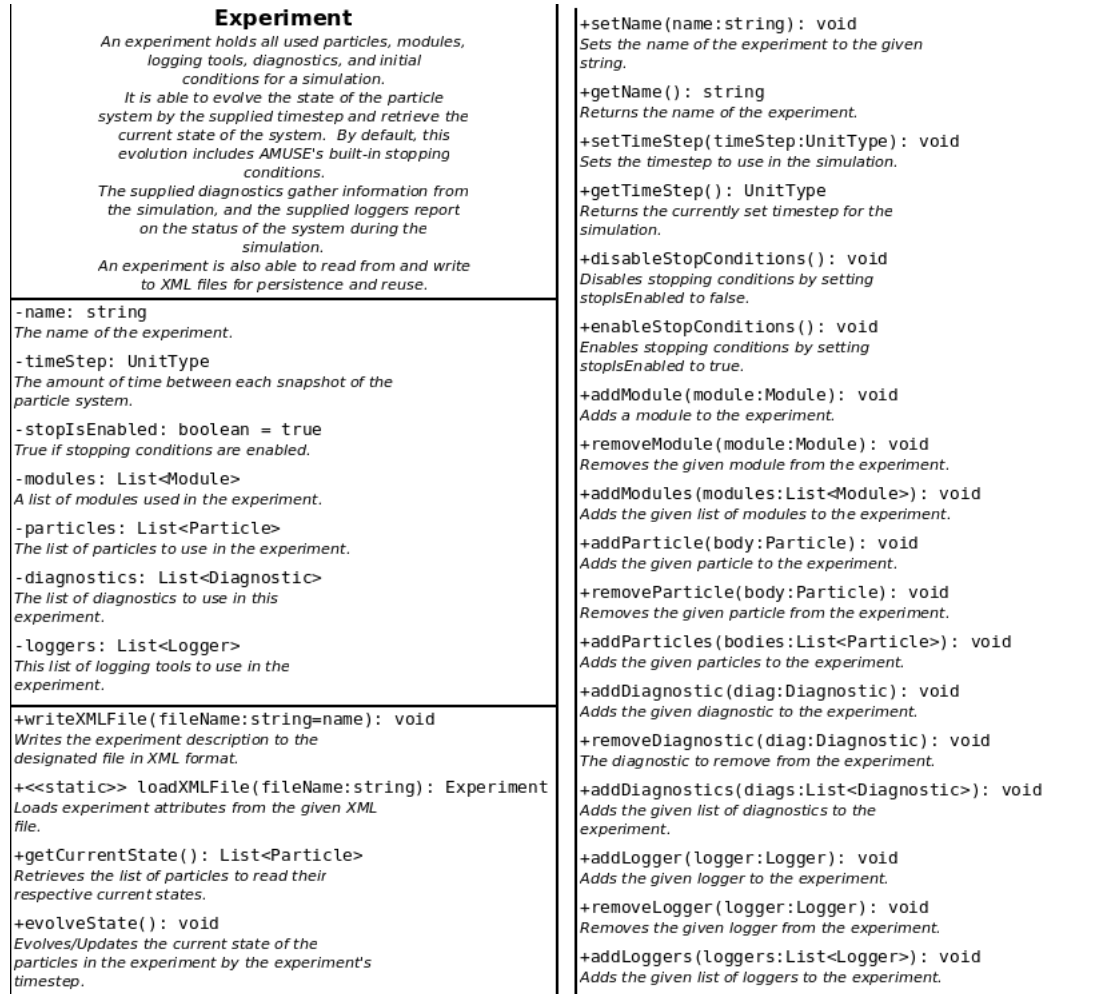


Figure 4.1: Experiment Class Diagram

4.1.1 Experiment Class

An instance of the Experiment class represents a single AMUSE experiment. It holds all particles (Satisfies requirement 2.3.3), modules, logging tools, diagnostics, and initial conditions for a simulation. These are described in more detail below. The experiment class can be viewed in Figure 4.1.

4.1.1.1 Experiment Attributes Experiments have the following attributes. The type UnitType is defined later in this document.

Name	Type	Description
name	string	The name of the experiment.
timeStep	UnitType	The amount of time between each snapshot of the simulation.
stopIsEnabled	boolean	Whether or not stopping conditions are enabled.
modules	List<Module>	The list of Modules used in the experiment.
particles	List<Particle>	The list of Particles to use in the experiment.
diagnostics	List<Diagnostic>	The list of Diagnostics to use in the experiment.
loggers	List<Logger>	The list of Loggers to use in the experiment.

4.1.1.2 Experiment Operations

<i>void writeXMLFile(string fileName)</i>		<i>static Experiment loadXMLFile(string fileName)</i>	
Input	A string containing the filename to write the Experiment XML file to.	Input	A string containing the filename to load the Experiment XML from.
Output	(none)	Output	The Experiment specified by the XML file.
Description	Writes the Experiment in XML format to the specified filename.	Description	Recreates an Experiment from its XML specification.
<i>List<Particle> getCurrentState()</i>		<i>void evolveState()</i>	
Input	(none)	Input	(none)
Output	The list of Particles in the Experiment.	Output	(none)
Description	Returns the list of Particles in the simulation with their updated attributes to read their current states.	Description	Evolves/Updates the current state of the particles in the Experiment by the Experiment's timeStep.

<i>void setName(string name)</i>		<i>string getName()</i>	
Input	A string containing the updated Experiment name.	Input	(none)
Output	(none)	Output	A string containing the name of the Experiment.
Description	Updates the Experiment's name to the given parameter.	Description	Returns the name of the Experiment.
<i>void setTimeStep(UnitType timeStep)</i>		<i>Units getTimeStep()</i>	
Input	A UnitType containing the updated Experiment timeStep.	Input	(none)
Output	(none)	Output	A UnitType containing the Experiment timeStep.
Description	Updates the Experiment's timeStep to the given parameter.	Description	Returns the timeStep of the Experiment.
<i>void disableStopConditions()</i>		<i>void enableStopConditions()</i>	
Input	(none)	Input	(none)
Output	(none)	Output	(none)
Description	Disables the Experiment's stopping conditions.	Description	Enables the Experiment's stopping conditions.
<i>void addModule(Module module)</i>		<i>void removeModule(Module module)</i>	
Input	A Module to add to the Experiment.	Input	A Module to remove from the Experiment.
Output	(none)	Output	(none)
Description	Adds the given Module to the Experiment.	Description	Removes the given Module from the Experiment.
<i>void addModules(List<Module> modules)</i>		<i>void addParticle(Particle body)</i>	
Input	A list of Modules to add to the Experiment.	Input	A Particle to add to the Experiment.
Output	(none)	Output	(none)
Description	Adds the given Modules to the Experiment.	Description	Adds the given Particle to the Experiment.

<i>void addParticles(List<Particle> bodies)</i>		<i>void removeParticle(Particle body)</i>	
Input	A list of Particles to add to the Experiment.	Input	A Particle to remove from the Experiment.
Output	(none)	Output	(none)
Description	Adds the given Particles to the Experiment.	Description	Removes the given Particle from the Experiment.
<i>void addDiagnostic(Diagnostic diag)</i>		<i>void addDiagnostics(List<Diagnostic> diags)</i>	
Input	A Diagnostic to add to the Experiment.	Input	A list of Diagnostics to add to the Experiment.
Output	(none)	Output	(none)
Description	Adds the given Diagnostic to the Experiment. Satisfies requirement 2.6.3.{1,2}.	Description	Adds the given Diagnostics to the Experiment.
<i>void removeDiagnostic(Diagnostic diag)</i>		<i>void addLogger(Logger logger)</i>	
Input	A Diagnostic to remove from the Experiment.	Input	A Logger to add to the Experiment.
Output	(none)	Output	(none)
Description	Removes the given Diagnostic from the Experiment.	Description	Adds the given Logger to the Experiment. Satisfies requirement 2.7.1.
<i>void addLoggers(List<Logger> loggers)</i>		<i>void removeLogger(Logger logger)</i>	
Input	A list of Loggers to add to the Experiment.	Input	A Logger to remove from the Experiment.
Output	(none)	Output	(none)
Description	Adds the given Loggers to the Experiment.	Description	Removes the given Logger from the Experiment. Satisfies requirement 2.7.2.

4.2 Particles

Particles are the elements that are being updated and the experiment is being run on. These particles can be anything as little as atoms to as large as star clusters. They are made up of mass, position, and velocity. (Satisfies requirements 2.3.1, 2.3.2)

4.2.1 Particle Class

Particle
-mass: Parameter -velocity: Parameter -position: Parameter
+getVelocity(): Parameter +setVelocity(v:List<double>,unit:Unit): boolean +getPosition(): Parameter +setPosition(pos:List<double>,unit:Unit): bool +getMass(): Parameter +setMass(mass:double,unit:Unit): bool

Figure 4.2: Particle Class Diagram

Particle class will be instantiated and used for storage of particle type items on the UI of GPUUnit. The Particle class in AMUSE will be used in the CLT. The particle class will be used primarily for data storage for experiment creation.

4.2.1.1 Particle Attributes

Name	Type	Description
mass	Parameter	The current mass of the particle. Will have units.
velocity	Parameter	The current mass of the particle. Will have units.
position	Parameter	The current position of the particle. Will have units.

4.2.1.2 Particle Methods

<i>Parameter getMass()</i>		<i>bool setMass(double mass, Unit unit)</i>	
Input	(none)	Input	The new value of the mass and the unit it is in.
Output	A parameter with the value and units of the current mass.	Output	Returns true if the mass is successfully set.
Description	Gets the current value of the mass.	Description	Sets the mass of the particle.
<i>Parameter getVelocity()</i>		<i>bool setVelocity(List<double> v, Unit unit)</i>	
Input	(none)	Input	The new vector value of the velocity and the unit it is in.
Output	A parameter with the value and units of the current velocity.	Output	Returns true if the velocity is successfully set.
Description	Gets the current value of the velocity.	Description	Sets the velocity of the particle.

<i>Parameter getPosition()</i>		<i>bool setPosition(List<double> pos, Unit unit)</i>	
Input	(none)	Input	The new vector value of the position and the unit it is in.
Output	A parameter with the value and units of the current position.	Output	Returns true if the position is successfully set.
Description	Gets the current value of the position.	Description	Sets the position of the particle.

4.3 Modules

This section defines an object model for the representation of and interaction with AMUSE modules, including several ancillary types utilized by the module implementation.

4.3.1 Module Class

Module
<pre> +name: string +description: string +domain: AstrophysicalDomain +codeName: string +codeLocation: string +isParallel: boolean +stoppingConditions: StoppingConditions +parameters: List<Parameter> +toXml(): string +fromXml(element:string): Module </pre>

Figure 4.4: Module class diagram

An instance of the Module class represents a particular AMUSE module, providing an interface between GPUnit and the AMUSE code. The members of the Module class are used by GPUnit to properly locate and initialize an AMUSE module, as well as display a module's details and configuration to the user in the graphical interface. (Satisfies requirements 2.2.1, 2.2.2, 2.2.3, 2.2.4, 2.2.5)

4.3.1.1 Module Attributes

Name	Type	Description
name	string	The full name of the module.
description	string	A description of the module's purpose (in other words, the calculations performed by the module).
domain	AstrophysicalDomain	The astrophysical domain into which the module has been sorted.
codeName	string	The name of the AMUSE class containing the module code.
codeLocation	string	The location of the AMUSE module code.
isParallel	boolean	Whether the module's calculations can be parallelized across multiple workers by MPI.
stoppingConditions	StoppingConditions	The condition(s) under which the module may stop executing prematurely.
parameters	List<Parameter>	The module parameters. These module-specific values may be modified prior to running the experiment in order to fine-tune the module's behavior.

4.3.1.2 Module Operations

<i>string toXml()</i>		<i>static Module fromXml(string element)</i>	
Input	(none)	Input	A string containing an XML representation of a Module.
Output	A string containing an XML representation of the Module.	Output	A Module whose property values are specified by the given XML element.
Description	Dumps the Module to an XML element.	Description	Recreates a Module from its XML specification.

4.3.2 Parameter Class




Figure 4.5: Parameter class diagram

Modules have parameters that may be modified by the user via the graphical interface. These parameters are typically specific to the module's domain or even to the individual module. The Parameter class is a generic class whose type parameter specifies the type of the parameter's value (e.g. integer or floating-point). A parameter may be a physical quantity, flag, or other value. (Satisfies Requirement 2.3.6)

4.3.2.1 Parameter Attributes

Name	Type	Description
T	Type	The data type of the parameter's value.
name	string	The brief, descriptive name of the parameter.
description	string	A description of the parameter's meaning and effects.
defaultValue	T	The parameter's default value.
minValue	T	A lower bound on the range of valid parameter values.
maxValue	T	An upper bound on the range of valid parameter values.
units	CompoundUnit	The physical unit(s) associated with the parameter's value.

4.3.2.2 Parameter Operations

<i>string toXml()</i>		<i>static Parameter<T> fromXml(string element)</i>	
Input	(none)	Input	A string containing an XML representation of a Parameter.
Output	A string containing an XML representation of the Parameter.	Output	A Parameter whose property values are specified by the given XML element.
Description	Dumps the Parameter to an XML element.	Description	Recreates a Parameter from its XML specification.

4.3.3 Unit Class




Figure 4.6: Unit class diagram

With AMUSE, physical quantities may be expressed using, and converted between, a number of different standard units. The Unit class, in conjunction with the CompoundUnit class, allows GPUUnit to present these units to the user for the purposes of description and selection.

4.3.3.1 Unit Attributes

Name	Type	Description
type	UnitType	The base type of astrophysical unit being represented.
prefix	SIPrefix	The SI prefix that modifies the base unit's order of magnitude.
exponent	integer	The exponent to be applied to the unit. For example, a typical unit of volume (such as cubic meters) would have an exponent of 3.

4.3.4 CompoundUnit Class




Figure 4.7: CompoundUnit class diagram

When performing physical calculations, units of measure may be combined in any number of ways. The CompoundUnit class provides the mechanism by which units may be presented to the user, regardless of whether they are simple or compound.

4.3.4.1 CompoundUnit Attributes

Name	Type	Description
description	string	A non-abbreviated textual description of the combined physical unit.
symbolicDescription	string	A shorthand textual description of the combined physical unit, using unit abbreviations.
units	List<Unit>	The list of simple units whose combination is represented by this compound unit.

4.3.5 UnitType Enumeration




Figure 4.8: UnitType type diagram

UnitType enumerates the myriad of base physical units supported by AMUSE. The list of members shown may not be exhaustive.

4.3.6 SIPrefix Enumeration




Figure 4.9: SIPrefix type diagram

When dealing with physical quantities, standard prefixes may be prepended to units to denote a given quantity's order of magnitude. SIPrefix enumerates the possible unit prefixes.

4.3.7 AstrophysicalDomain Enumeration




Figure 4.10: AstrophysicalDomain type diagram

The modules included with AMUSE have categorized according to their domain. A module's domain indicates the quantities it deals with as well as common operations that may be called on the module. The **AstrophysicalDomain** type enumerates the several domains that have been specified by AMUSE.

4.3.8 StoppingConditions Enumeration

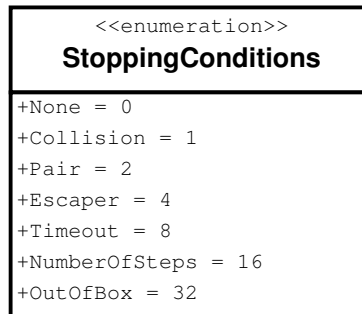


Figure 4.11: StoppingConditions type diagram

StoppingConditions is a set of bit flags enumerating the conditions under which a module may return from execution before completing its assigned calculations.

4.4 Diagnostics

Diagnostics have access to the state of the experiment at the end of each time step. Default diagnostics will be provided, e.g. printing the state to a file at a specified time step or creating a stellar mass histogram. Users can create custom diagnostic plug-ins. (Satisfies requirement 2.6)

4.4.1 Object Model

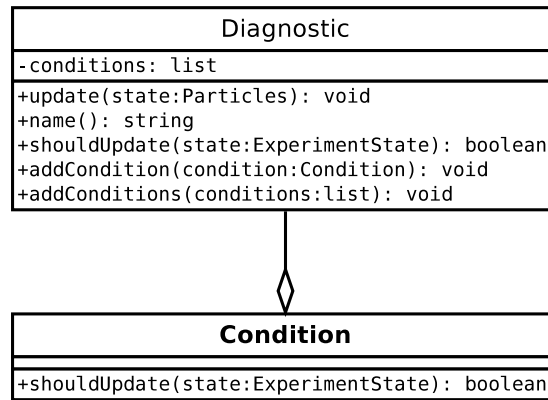


Figure 4.12: Diagnostics Object Model

4.4.2 Diagnostic Class

The Diagnostic class is an abstract class that is the base of both built-in and custom diagnostics tools. All custom diagnostic scripts must inherit from this class. Subclasses must implement the update and name operations. The Diagnostic class also contains a list of conditions to check if the diagnostic needs to be updated.

4.4.3 Experiment Manager Class

The Experiment Manager class has the responsibility to manage as well as update the active diagnostics.

4.4.4 Condition Class

The Condition class represents the conditions that need to be satisfied for a diagnostic to be updated.

4.5 Logging

Logs are helpful when the experiment must be restarted along with its output data if the experiment has an error or needs to be stopped for any other reason. Logging could be enabled or disabled by the user. The output destination for the logs could be specified. The specified destination for the output can be one among the following: the console, a file, or a window in the GUI. (Satisfies requirement 2.7)

4.5.1 Object Model

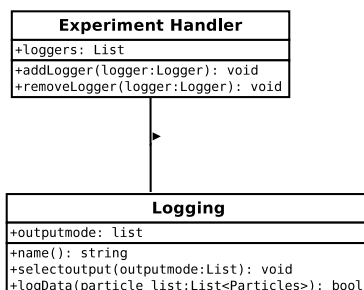


Figure 4.13: Logging Object Model

4.5.2 Logging Class

The Logging class takes care of selecting the required output option as the console, file or the GUI window.

5 Command Line Tool

5.1 Overview

The command line tool (CLT) is used for the running of an experiment file. The CLT begins with parsing the experiment file. After parsing the experiment file the CLT takes the initialization parameters and loads them using the AMUSE framework. The CLT then takes the modules specified, searches through loaded AMUSE modules and initializes the modules. The CLT links the modules together to allow interaction and updating of the current state of the system using the AMUSE channels. The CLT then loads the Logging and Data Analysis scripts specified by the experiment file, linking it to the particle array. The CLT then runs through the time specified in the experiment file, evolving all of the modules loaded at each time step. It will run the data analysis scripts as well as the logging on the intervals specified. After the model is done evolving it will output using whatever data analysis script listed, in the file format specified.

5.2 Using the CLT

The command line tool will primarily be used by the GUI to send out jobs and run them to make the user experience much more enjoyable and at ease. The command line tool though can be used with a basic experiment file with some on the fly modifications by using some flags.(Satisfies requirement 2.4.1)

5.2.1 CLT Flags

Flag	<Value Being Passed>	Description
-help	None	Displays a useful help prompt with the list of CLT flags.
-f	Experiment File Location	Loads the experiment file prompted by the filename.
-n	Number of Particles	Changes the number of particles being passed in the experiment file.
-t	End Time	Changes the max time allowed in the simulation.
-dt	Timestep	Changes the timestep in which dt moves forward.
-r	Radius	Changes the radius scaling factor.

5.3 Running The Experiment

5.3.1 Simulation Overview

A Simulation is the goal of both GPUnit and AMUSE, which makes it important to understand and break it into key parts. The Simulation can be looked at as three sections: Simulation Initialization, Running Simulation, Post Simulation. This section will go into more detail about all three of these parts and what happens in each.




Figure 5.1: Flow Chart of Typical AMUSE Simulation

5.3.1.1 Simulation Initialization Simulation Initialization is used for preparation of the particles and modules to be able to model the system. The following will break down the initialization into the main steps.

5.3.1.1.1 Instantiate Particles In the initialization phase of the experiment it is important to select proper initial conditions. The main initial conditions to set is the mass distribution and starting locations of the particles. There are built in AMUSE classes that are used to set these. The CLT while reading in the experiment file will find the initial condition files by name. (Satisfies Requirements 2.3.4)

5.3.1.1.2 Create Conversion Object Most astronomical simulations use a scaling factor so you can do the numerical integration quicker. A conversion object is

a class in AMUSE that you instantiate by passing in the total mass and the average radius of the system. You will then pass this conversion object into any modules you will be using in the future so the module knows how to translate between this smaller number format to an actual real data value.

5.3.1.1.3 Module Creation The CLT needs to instantiate the AMUSE class for the specific module it is working on from the location specified in the module class. It needs to instantiate the module being passed in the conversion object so the module knows what units it is working in and how to convert its units. The module then sets its list of particles to the master list of particles. After it does that it needs to run its initial calculations needed for evolving the module.

5.3.1.1.4 Instantiate Channels Channels are used to be able to communicate data to and from the modules that may be on a different system. On every time step you need to synchronize the master list of particles to every modules separate list of nodes. Afterwards you need to re-sync the modules nodes to the master node. This allows the modules to have the same data to move forward in the simulation. (Satisfies Requirement 2.5.1, 2.5.2, 2.5.2.1)

5.3.1.2 Running Simulation Running the system is the area inside of the while loop. This is where the modules are run, some logging is done, some diagnostics and increment the time scale. (Satisfies Requirement 2.4.2)

5.3.1.2.1 Evolve Modules The evolution of the Module is core of running a simulation. The evolution of the module is done by running the AMUSE module's `evolve_model` function with passing in the current time.

5.3.1.2.2 Synchronize Particles Synchronize particles is required due to the fact every module has a local list of particles to allow parallelization. These particles are updated, yet only one set of dynamics or stellar evolution will be run at a time which prevents any overlapping of data modification. To synchronize the particles you run all the channels created to synchronize all the data to the master list of particles. After the master list is up to date, you then synchronize using all the channels to the modules. This leaves all the modules' list of particles with the latest version of particle data.

5.3.1.2.3 Run Diagnostics The function should loop through all the diagnostic scripts and if the diagnostics should use the `shouldUpdate` function with the current state of the experiment to check if it should run. If the diagnostic should run, it then runs the update method passing in the master list of particles.(Figure 4.12)

5.3.1.2.4 Run Logging The function will loop through all the logging scripts and run any real time logging scripts that need to be run. The real time logging scripts are denoted by name and run their logData method passing in the list of particles.(Figure 4.13)

5.3.1.2.5 Increment Time The time should increment using the specified timestep (Figure 4.1). The stopping of the experiment is signaled when the stopping conditions are met. (Satisfies Requirement 2.4.3.1)

5.3.1.3 Post Simulation

5.3.1.3.1 Run Closing Diagnostics Scripts The function should loop through all the diagnostic scripts and if the diagnostics should use the shouldUpdate function with the current state of the experiment to check if it should run. If the diagnostic should run, it then runs the update method passing in the master list of particles.(Figure 4.12)

5.3.1.3.2 Run Closing Logging Scripts The function will loop through all the logging scripts and run any post simulation logging scripts that need to be run. The post simulation logging scripts are denoted by name and run their logData method passing in the list of particles.(Figure 4.13)

5.3.1.3.3 Stop Modules The modules should be cleaned up after being run. This involves if running in parallel sending a kill command via MPI. We will run the modules'(AMUSE class) stop() command.

6 Networking

6.1 Overview

The networking layer is responsible for communication between the interface and the cluster status daemons running on the cluster nodes. The user interface and console can send messages through the network which will be distributed via IP multicast to any cluster nodes listening to a selected multicast group. Cluster nodes will be able to send replies to the sender containing information such as health stats and the number of experiments running on that node. The UI and command line tools will use IPC (pipes etc...) to communicate with local instances of the control program to send network packets to the cluster nodes.

6.2 Network Software Components

6.2.1 Control Instance

A control instance is a program that will run in the background on any machine wishing to control experiments on a cluster. The core scripts (command line tools)

and UI will both communicate with this instance to perform any network access.

6.2.2 Node Instance

Node instances are processes that run on each computing node in the cluster. The node instance tracks health statistics and running experiments and can send these pieces of information as direct (unicast) replies to queries received through the multicast group.

6.3 Object Model

6.3.1 Overview

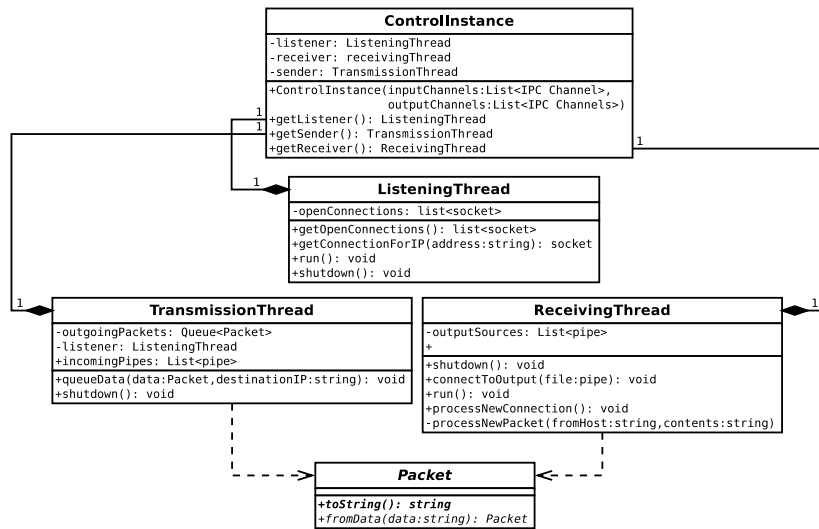


Figure 6.1: Network Object Model

6.3.2 Control Instance




Figure 6.2: Control Instance Class Diagram

6.3.3 Node Instance

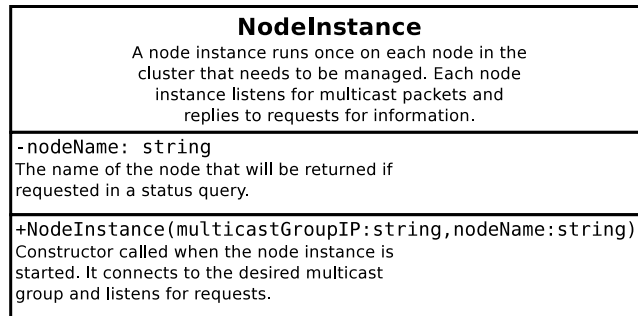


Figure 6.3: Node Instance Class Diagram

6.3.4 TransmissionThread

The transmission thread (Figure 6.4) will process outgoing packets queued for transmission by external sources such as the UI.

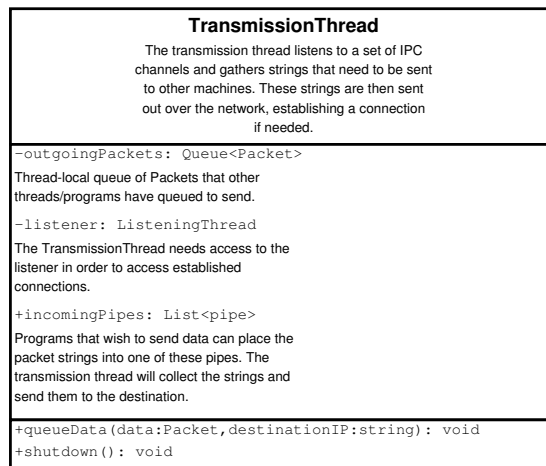


Figure 6.4: Transmission Thread Class Diagram

6.3.5 ReceivingThread

The receiving thread (Figure 6.5) will read incoming packets from established connections with node instances.

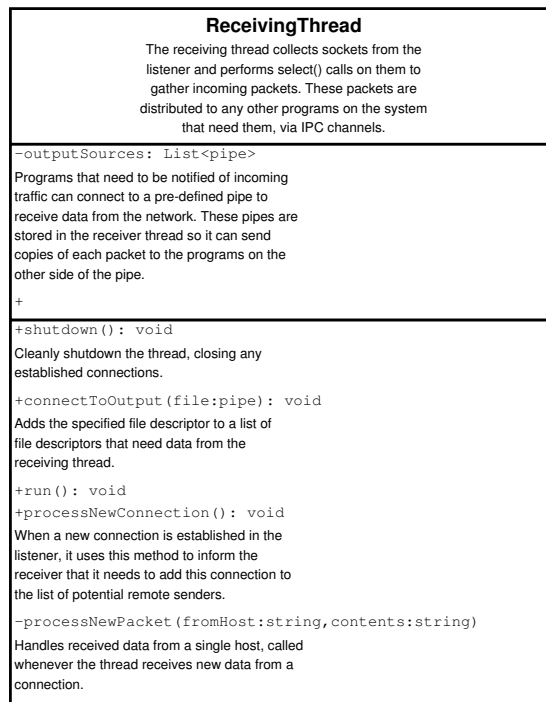


Figure 6.5: Receiving Thread Class Diagram

6.3.6 ListeningThread

The listening thread (Figure 6.6) will listen for and accept new incoming connections from node instances.

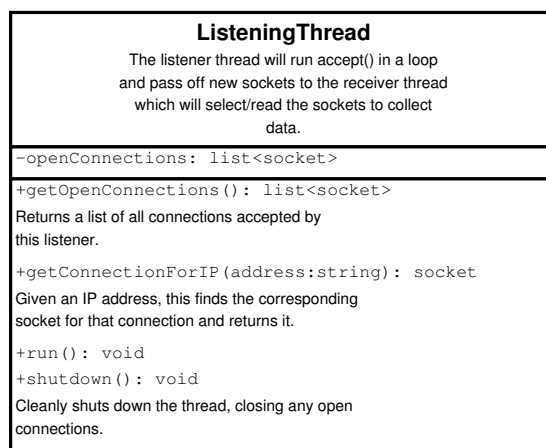


Figure 6.6: Listening Thread Class Diagram

6.3.7 Packet Objects

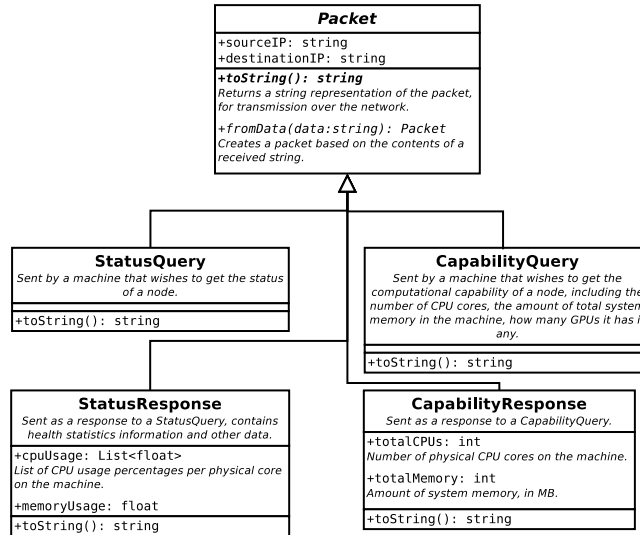


Figure 6.7: Packet Classes

7 Data

7.1 Experiment Specification

The following is a list of XML specifications for experiments. Experiments are saved and loaded via this format. Note that not all specified attributes will always be present because not all attributes are relevant to every experiment. See Figure 7.1 for an example XML experiment file. (Satisfies requirement 2.3.3, 2.3.5)

```

<experiment name="myExperiment">
  <time units="Myr" start="0" step="300" end="900"/>
  <module type="BHtree">
    <param name="use_self_gravity" value="0"/>
  </module>
  <particle radius="6371|km" mass="5.9736e4|kg" position="[1,0,0]|AU" velocity="[0.0,29783,0.0]|ms"/>
  <particle radius="1|RSun" mass="1|MSun" position="[0,0,0]|m" velocity="[0.0,0.0,0.0]|ms"/>
  <diagnostic name="massStats" file="massDiag.py"/>
  <diagnostic name="velocityStats" file="vDiag.py"/>
  <logger name="stateUpdates" file="stateUpdates.py"/>
  <logger name="timeSteps" file="timeSteps.py"/>
  <stopCondition type="COLLISION_DETECTION"/>
</experiment>

```

Figure 7.1: Example XML Experiment File

7.1.1 Experiment

Experiment tags are the first and last tags in an experiment specification because they encapsulate all experiment attributes.

Attribute	Description
name	The name of the experiment

7.1.2 Time

Time tags specify the time parameters of the experiment.

Attribute	Description
units	The time units for the experiment to use, defined by AMUSE
start	The start time of the experiment
step	The amount of time between simulation snapshots
end	The time to end the experiment

7.1.3 Module

Module tags define the modules used in the experiment. Each module has child param tags to define initial conditions. These initial conditions are defined by AMUSE.

Attribute	Description
name	The name of the module to use

7.1.4 Param

Param tags define the parameters for their parent model. Each defines a model's parameter value for use in the experiment. All parameters, default values, units, and descriptions are defined by AMUSE.

Attribute	Description
name	The name of the parameter to initialize
value	The value to give the parameter

7.1.5 Particle

Particle tags define the particles used in the experiment. All references to units are those defined by AMUSE.

Attribute	Description
mass	The mass of the particle with units
radius	The radius of the particle with units
position	The position of the particle in [X,Y,Z] format with units
velocity	The velocity of the particle in [X,Y,Z] format with units
luminosity	The luminosity of the particle with units
temperature	The temperature of the particle with units
age	The age of the particle with units
stellarType	The stellar type of the particle, defined by AMUSE

7.1.6 Diagnostic

Diagnostic tags define the diagnostic scripts used in the experiment.

Attribute	Description
name	The name of the diagnostic for reference purposes
file	The filename of the diagnostic script

7.1.7 Logger

Logger tags define the logging scripts used in the experiment.

Attribute	Description
name	The name of the logger for reference purposes
file	The filename of the logging script

7.1.8 Stopping Condition

Stopping Condition tags define the stopping conditions used in the experiment. All stopping conditions are defined by AMUSE.(Satisfies Requirement 2.4.3, 2.4.3.2)

Attribute	Description
type	The type of stopping condition to use

7.2 Network Packets

The following is a list of packet data content specifications. Data is described in terms of types to ease implementation, however the information is sent over the network as a string to avoid any machine-specific byte ordering issues. Fields in the packet are separated by a "|" character. If this character is to appear as text inside a packet for any reason, it must be escaped as "\\|"

7.2.1 Packet Header

The packet header precedes data in all packets.

PACKET_TYPE	LENGTH	SOURCE_IP	DEST_IP
int	long	string	string

7.2.2 Status Query Packet

The status request may contain flags requesting additional data from the node beyond the elements specified here. If the recipient understands the flags, they will fill the ADDITIONAL_DATA field with the appropriate response.

HEADER	ADDITIONAL_REQUEST_FLAGS
-	string

7.2.3 Status Response Packet

HEADER	CPU_USAGE	MEMORY_USAGE	SIMS_RUNNING	ADDITIONAL_DATA
-	float	float	string	string

7.2.4 Capability Query Packet

The additional data protocol is the same here as in the status request above.

HEADER	ADDITIONAL_REQUEST_FLAGS
-	string

7.2.5 Capability Response Packet

HEADER	NUM_CPUS	MAX_MEMORY	NUM_GPUS	ADDITIONAL_DATA
-	int	long	int	string

8 Appendix

8.1 Requirements Traceability Table

Requirement	Section in SDD
2.2.1	4.3.1
2.2.2	4.3.1
2.2.3	4.3.1
2.2.4	4.3.1
2.2.5	4.3.1
2.3.1	4.2
2.3.2	4.2
2.3.3	4.1
2.3.4	5.3.1.1.1
2.3.5	7.1
2.3.6	4.3.2
2.4.1	5.2
2.4.2	5.3.1.2
2.4.3	5.3.1.2.5
2.4.4	5.3.1.2.5
2.5.1	5.3.1.1.4
2.5.2	5.3.1.1.4
2.6.1	4.4
2.6.2	4.4
2.6.3	4.4
2.6.4	4.4
2.6.5	4.4
2.6.6	4.4
2.7.1	4.5
2.7.2	4.5
2.7.3	4.5
2.8.1	3.3
2.8.2	3.3

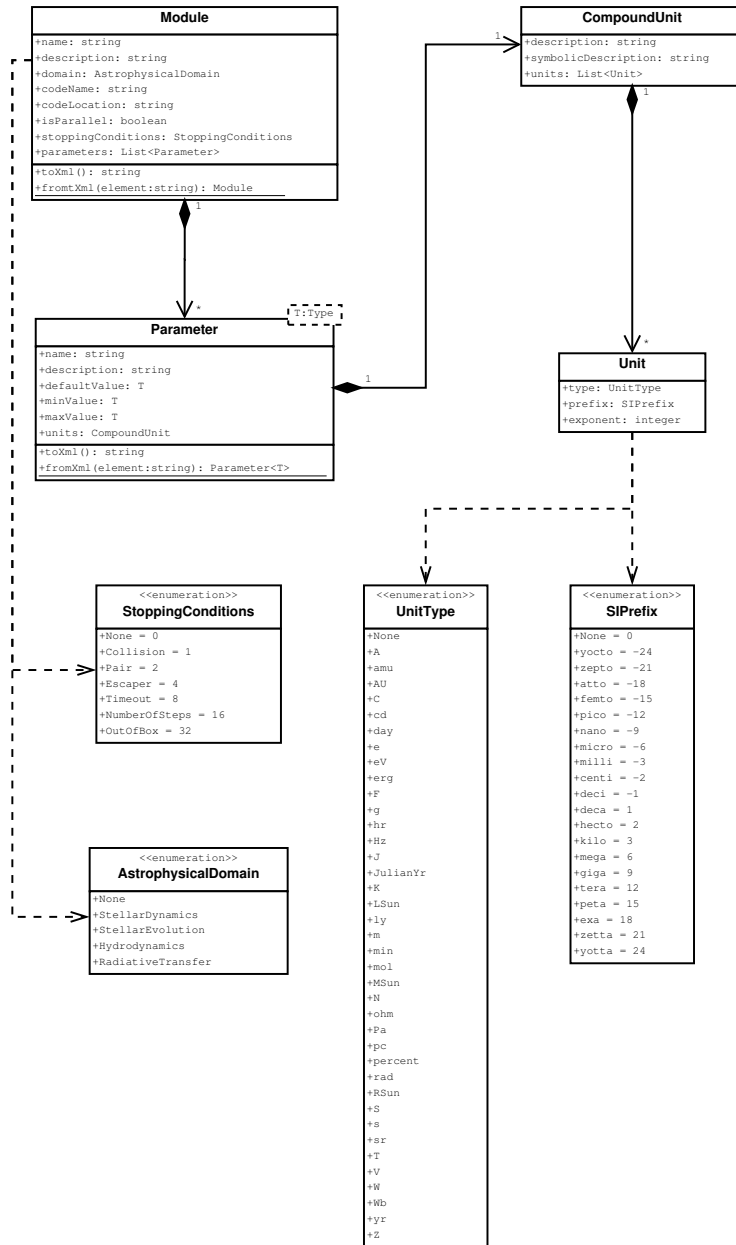


Figure 4.3: Relationships between Module and supporting classes